

□

# CS143A

## Principles on Operating Systems

### Discussion 02:

### OS Interfaces

Instructor: Prof. Anton Burtsev

TA: Saehanseul Yi (Hans)

Oct 16, 2019 **Noon**

# About me

- Link for all office hours/discussion:  
<https://uci.zoom.us/j/93369206818>

- Teaching staff office hours:

Hari: *Mon 12:00 PST*

Zhaofeng Li: *Tue 12:00 PST*

Deep: *Wed 9:00 AM PST*

**Hans: *Thu 12:00 PST***

Se-Min Lim: *Fri 9:00 PST*

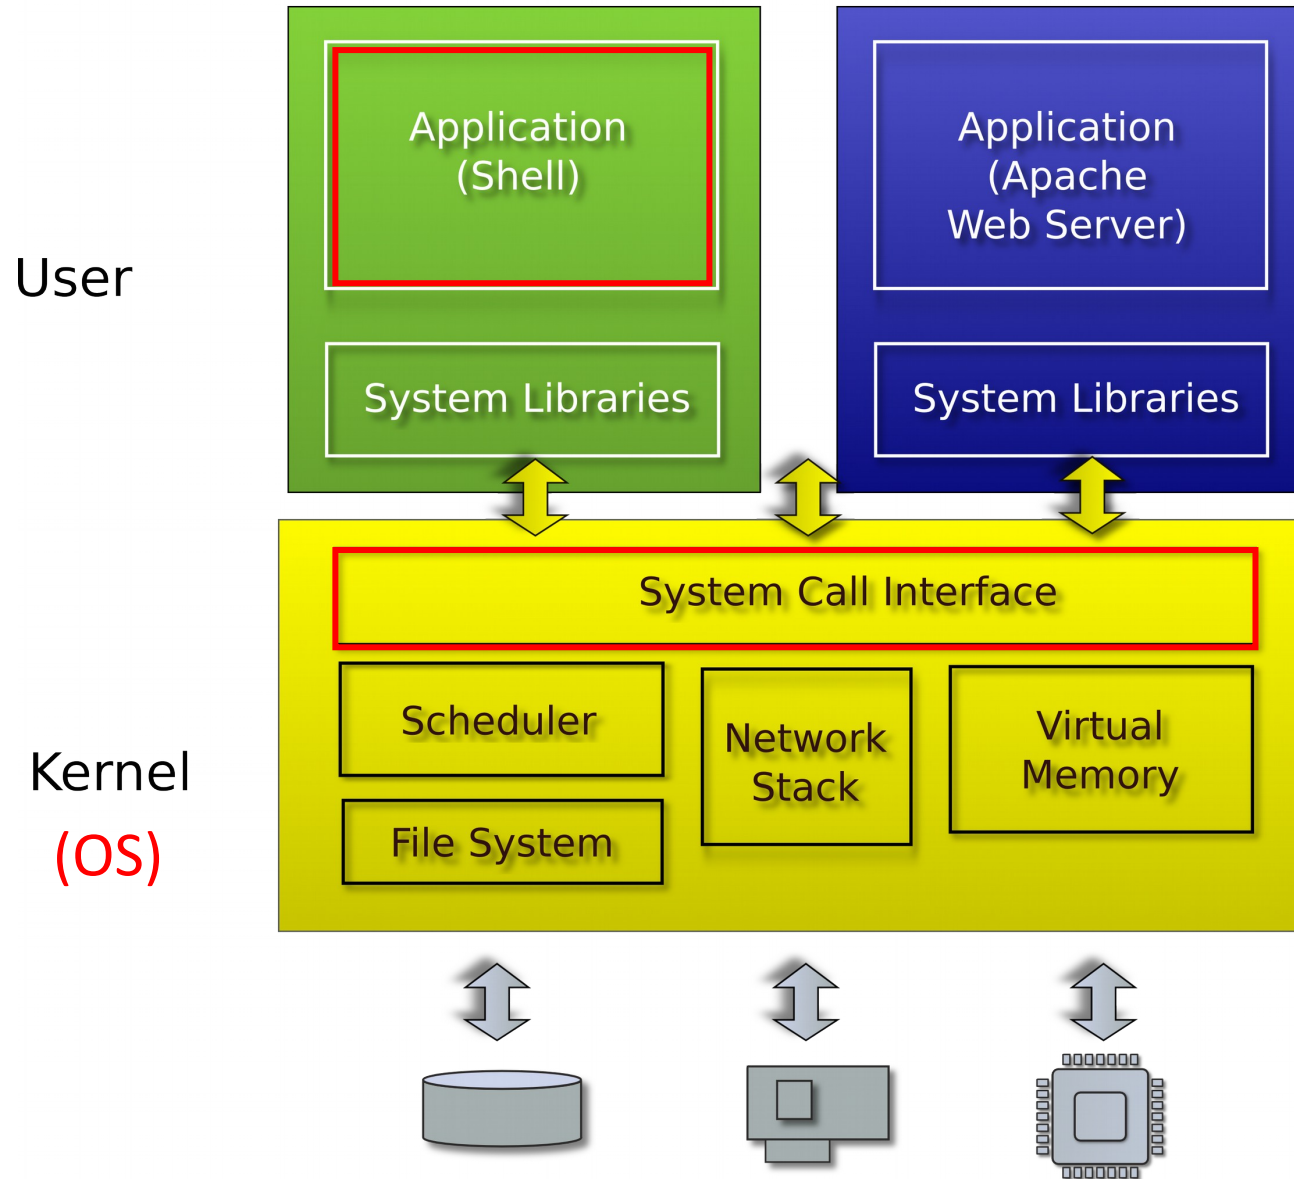
# Motivating example: redirection

- Best example for explaining `pipe()`, `fork()`, and `exec()`
- Program output -> *stdout* (default: screen)
- `|` (pipe *operator*): send outputs to somewhere else

```
$  
$ ls  
a.out  b.out  asdfasdf  
$
```

```
$  
$ ls | grep asdf  
asdfasdf  
$
```

# Typical UNIX OS



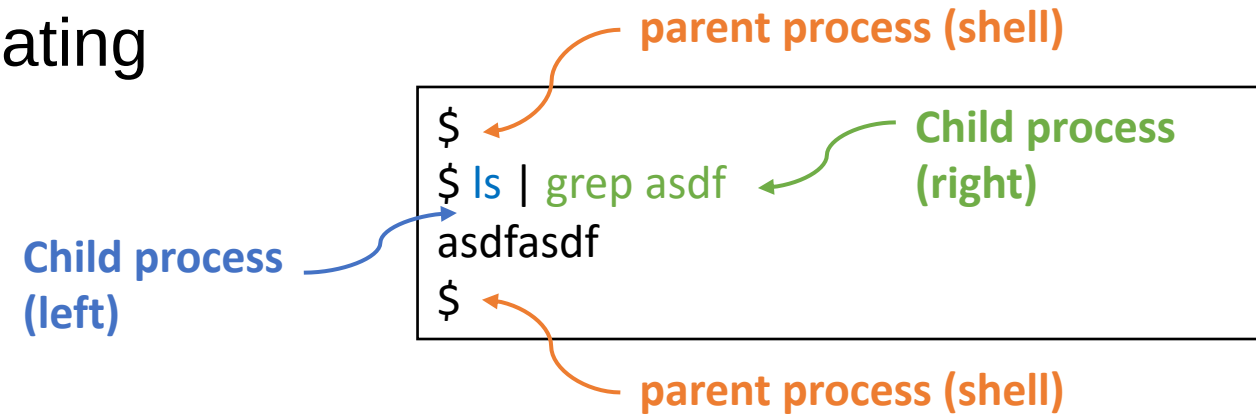
System calls are the interface of the OS

# But what is shell?

- Normal process
  - Kernel starts it for each user that logs in into the system
  - In xv6 shell is created after the kernel boots
- Shell interacts with the kernel through system calls
  - E.g., starts other processes

# System calls, interface for...


- Processes
  - **Creating**, exiting, waiting, terminating
- Memory
  - Allocation, deallocation
- Files and folders
  - Opening, reading, writing, closing
- Inter-process communication
  - **Pipes**



# Wait... stdin? stdout?

(standard input, standard output)

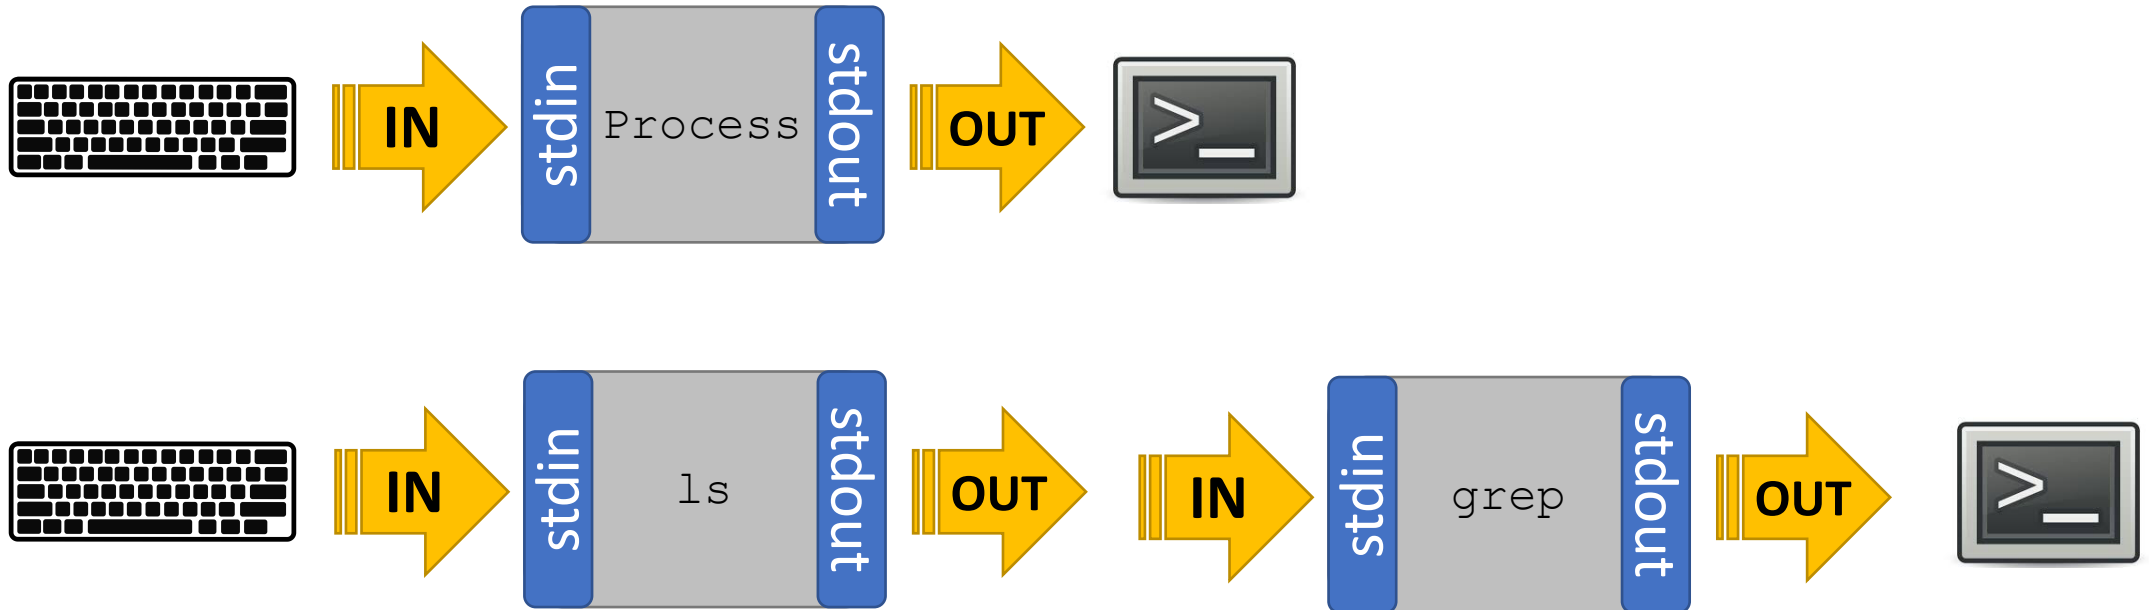
```
$  
$ ls | grep asdf  
asdfasdf  
$
```



# Wait... stdin? stdout?

(standard input, standard output)

```
$  
$ ls | grep asdf  
asdfasdf  
$
```






# Wait... stdin? stdout?

(standard input, standard output)

```
$  
$ ls | grep asdf  
asdfasdf  
$
```




# Wait... stdin? stdout?

(standard input, standard output)

```
$  
$ ls | grep asdf  
asdfasdf  
$
```

- stdin(0), stdout(1), and stderr(2) are file descriptors(**just an integer** in user-program)

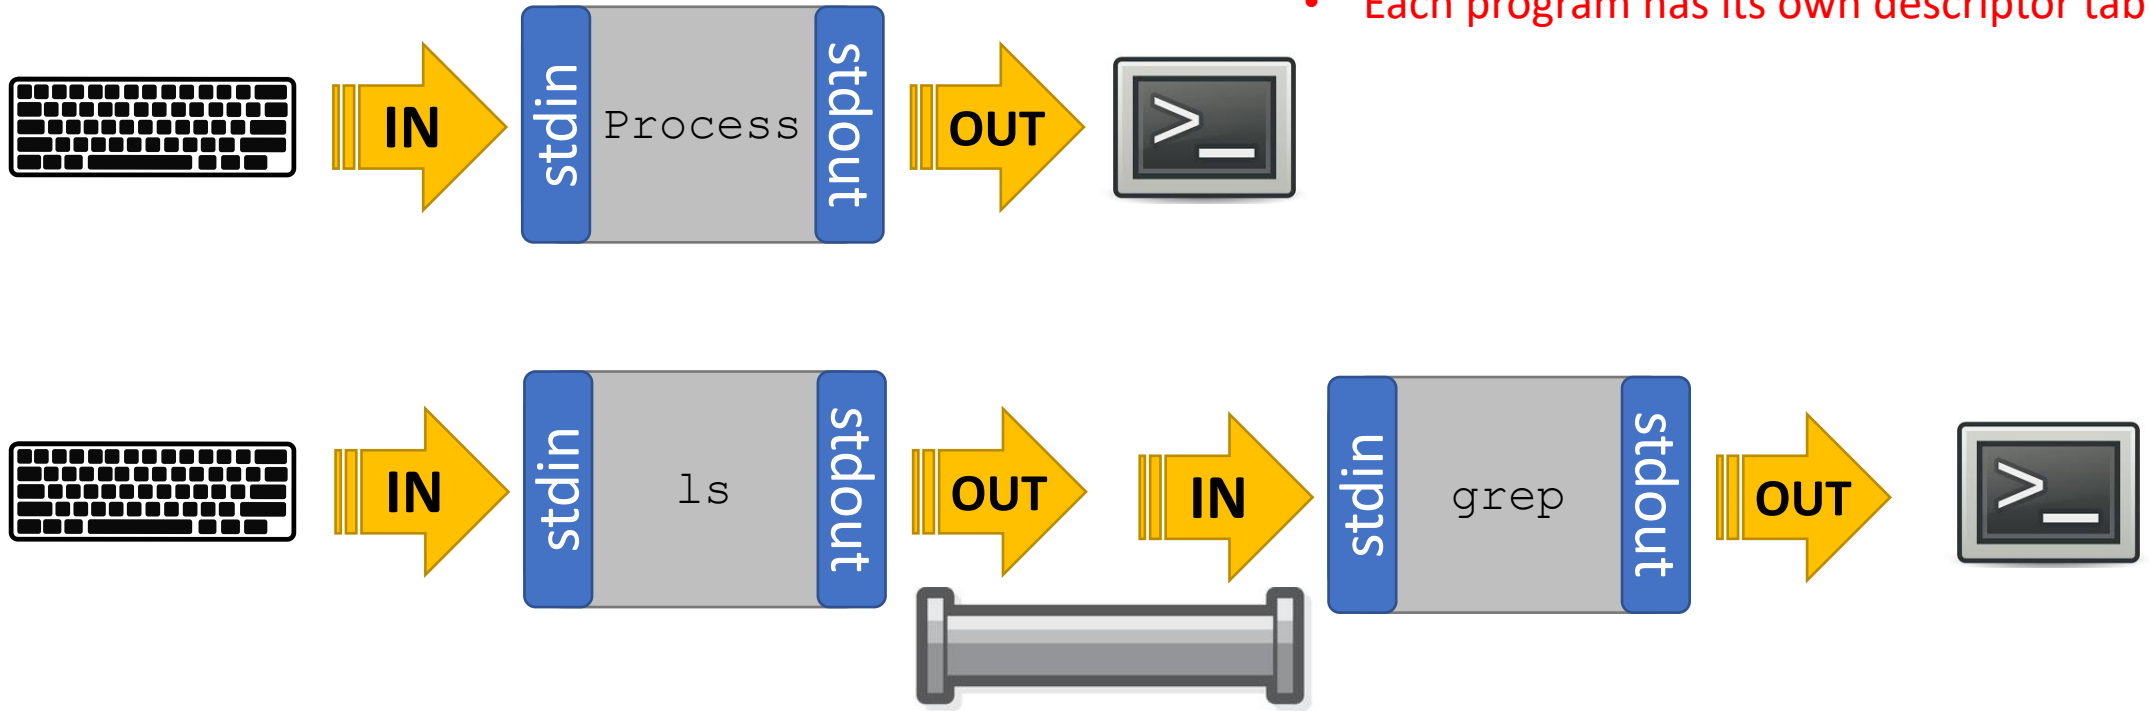


# Wait... stdin? stdout?

(standard input, standard output)

```
$  
$ ls | grep asdf  
asdfasdf  
$
```

- stdin(0), stdout(1), and stderr(2) are file descriptors(i.e. **just an integer** in user-program)
- Each program has its own descriptor table

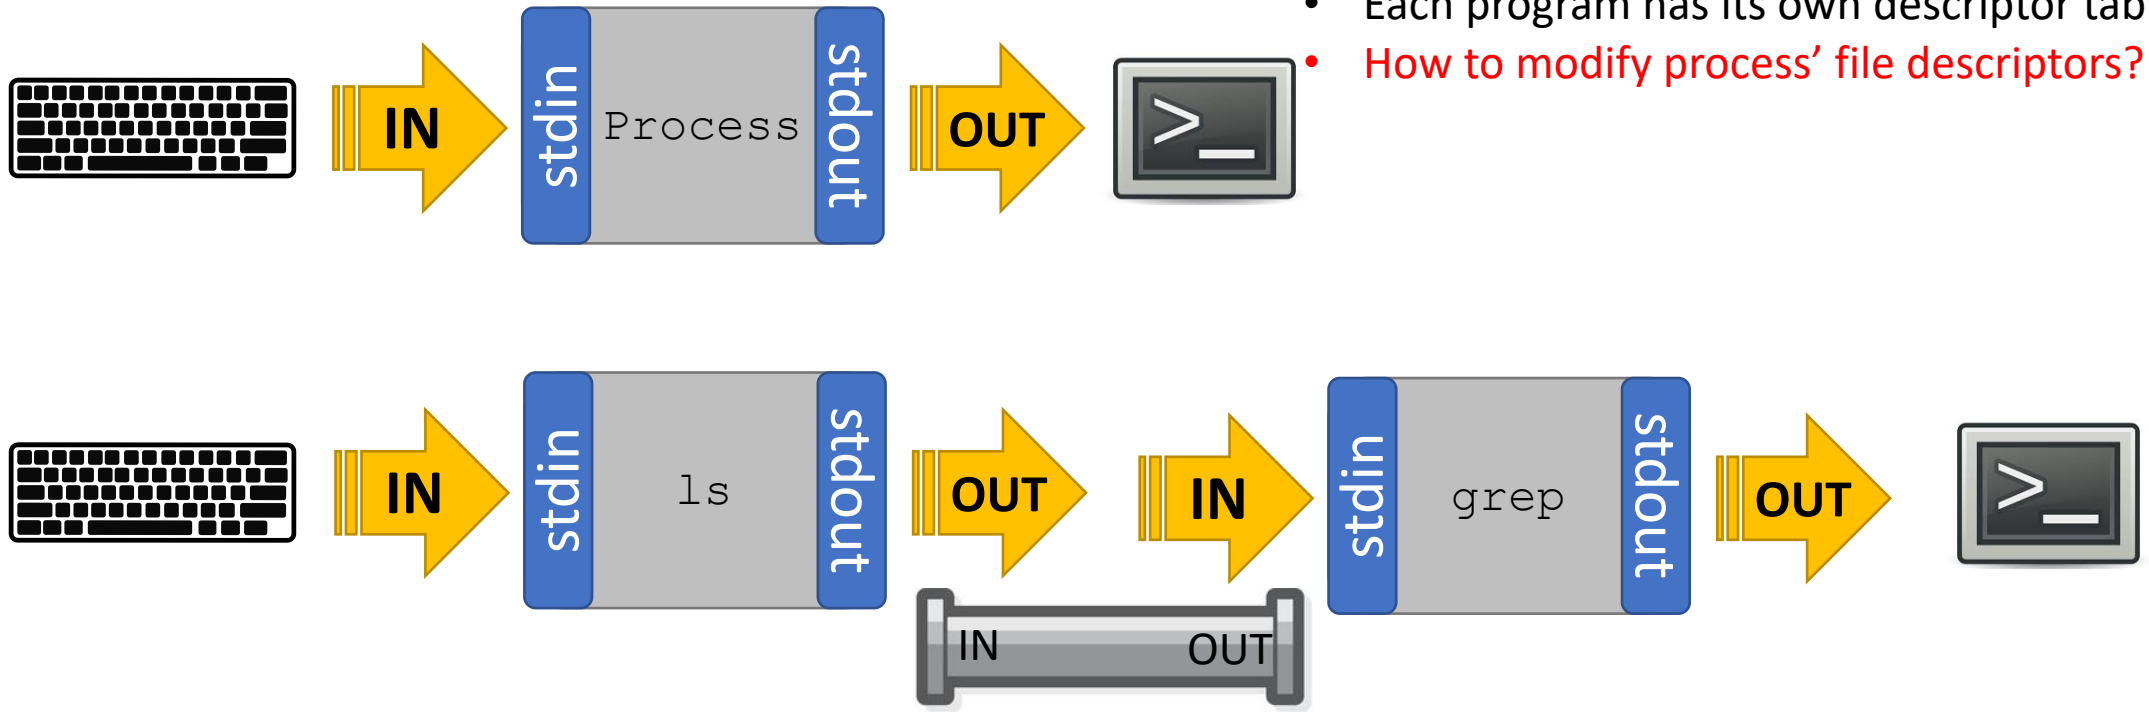


# Wait... stdin? stdout?

(standard input, standard output)

```
$  
$ ls | grep asdf  
asdfasdf  
$
```

- stdin(0), stdout(1), and stderr(2) are file descriptors(i.e. **just an integer** in user-program)
- Each program has its own descriptor table
- **How to modify process' file descriptors?**

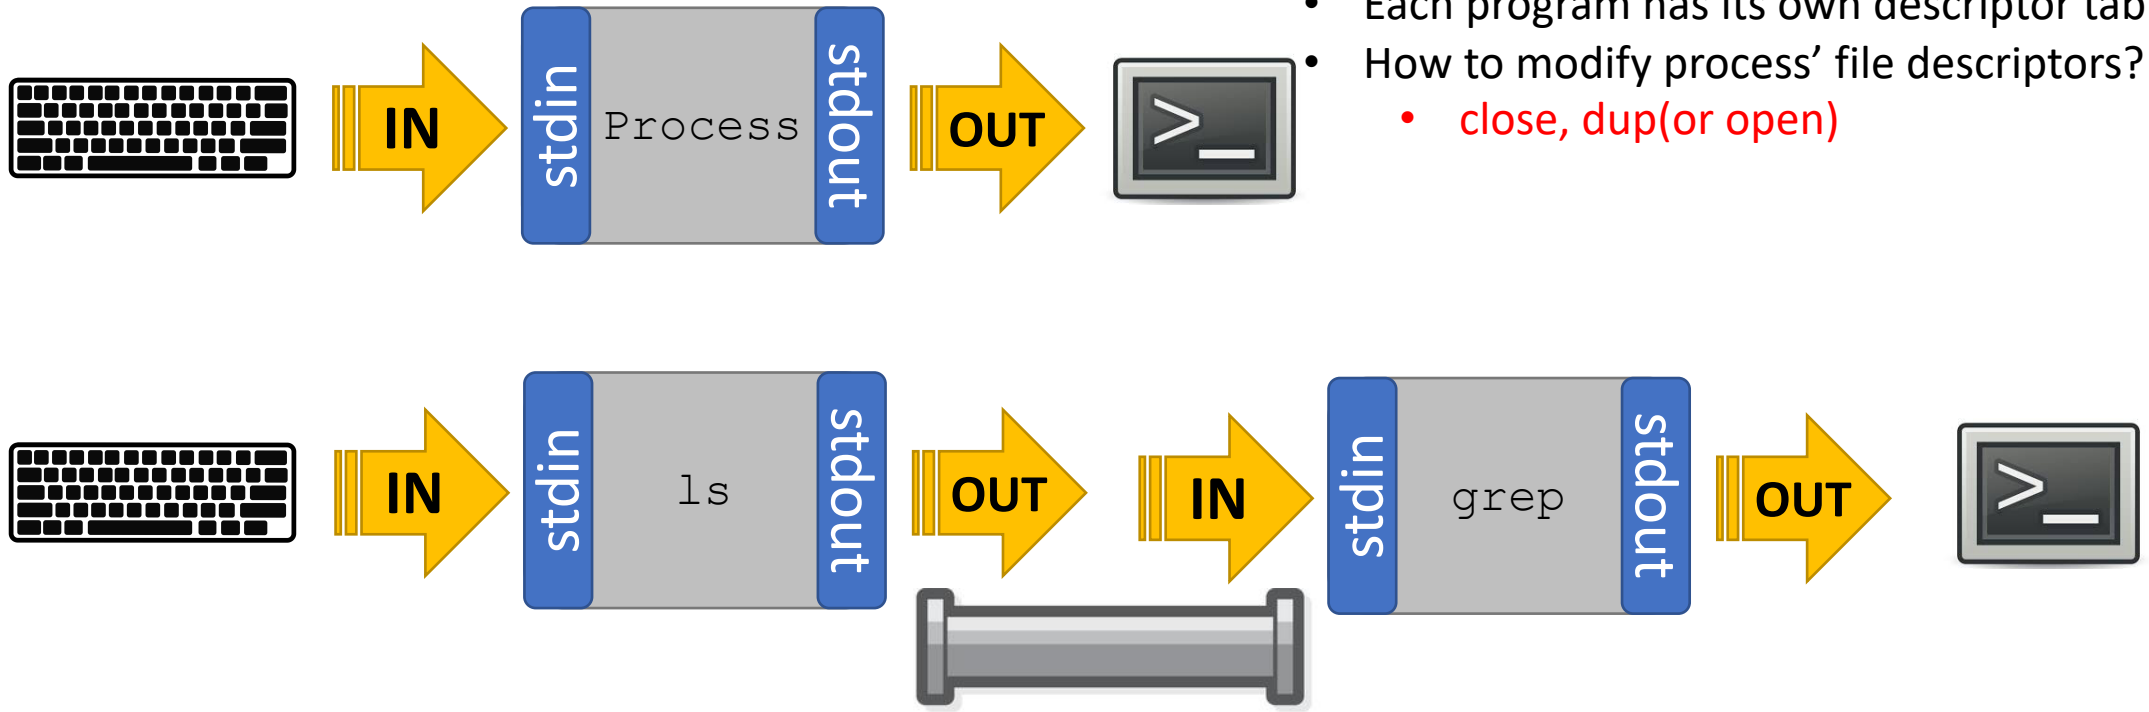


# Wait... stdin? stdout?

(standard input, standard output)

```
$  
$ ls | grep asdf  
asdfasdf  
$
```



- stdin(0), stdout(1), and stderr(2) are file descriptors(**just an integer** in user-program)
- Each program has its own descriptor table
- How to modify process' file descriptors?
  - **close, dup(or open)**



# Wait... stdin? stdout?

(standard input, standard output)

```
$  
$ ls | grep asdf  
asdfasdf  
$
```




- stdin(0), stdout(1), and stderr(2) are file descriptors(**just an integer** in user-program)
- Each program has its own descriptor table
- How to modify process' file descriptors?
  - close, dup(or open)
- What we need to do:  
close appropriate descriptors for each process  
and set the appropriate descriptor by copying

# Wait... stdin? stdout?


(standard input, standard output)

```
$  
$ ls | grep asdf  
asdfasdf  
$
```



- `stdin(0)`, `stdout(1)`, and `stderr(2)` are file descriptors(**just an integer** in user-program)
- Each program has its own descriptor table
- How to modify process' file descriptors?
  - `close`, `dup(or open)`

What we need to do:  
close appropriate descriptors for each process  
and set the appropriate descriptor by copying



POSIX.1-2001

`pipe()` creates a pair of file descriptors, pointing to a pipe inode, and places them in the array pointed to by `filedes`. `filedes[0]` is for reading, `filedes[1]` is for writing.

**pipe is uni-directional**

# pipe() and fork()

-----Point 0-----

case PIPE:

pcmd = (struct pipecmd\*)cmd;

if(**pipe**(p) < 0)  
 panic("pipe");

-----Point A-----

if(**fork1**() == 0){  
 close(1);  
 dup(p[1]);  
 close(p[0]);  
 close(p[1]);

-----Point B-----

runcmd(pcmd->left);  
}

if(**fork1**() == 0){  
 close(0);  
 dup(p[0]);  
 close(p[0]);  
 close(p[1]);  
 runcmd(pcmd->right);  
}

close(p[0]);  
close(p[1]);

-----Point C-----

wait();  
wait();  
break;

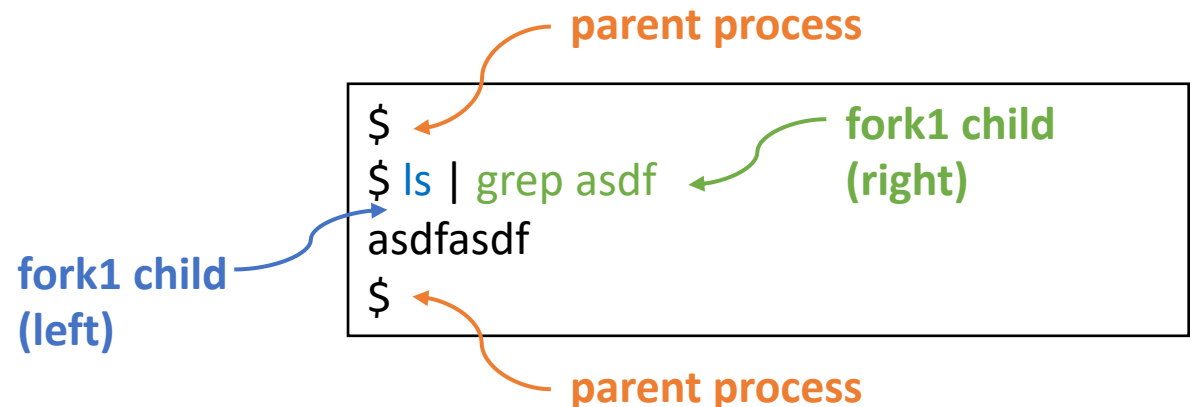
**int pipe(int pipefd[2]);**

Create a pipe & assign each end to pipefd

**pid\_t fork(void);**

Copy the current process (parent)

Returns the PID of the child (parent)  
or 0 (child)





# pipe() and fork()

※ Throughout the example, stderr is always connected to the screen. Omitted for simplicity as well as p[0] and p[1] to the parent process

-----Point 0-----


```
case PIPE:
pcmd = (struct pipecmd*)cmd;
if(pipe(p) < 0)
    panic("pipe");
```

-----Point A-----

```
if(fork1() == 0){
    close(1);
    dup(p[1]);
    close(p[0]);
    close(p[1]);
```


-----Point B-----

```
    runcmd(pcmd->left);
}
```



# pipe() and fork()

```
-----Point 0-----  
case PIPE:  
pcmd = (struct pipecmd*)cmd;  
if(pipe(p) < 0)      int p[2]  
    panic("pipe");  
-----Point A-----  
if(fork1() == 0){  
    close(1);  
    dup(p[1]);  
    close(p[0]);  
    close(p[1]);  
-----Point B-----  
    runcmd(pcmd->left);  
}
```



# pipe() and fork()

-----Point 0-----

```
case PIPE:
pcmd = (struct pipecmd*)cmd;
if(pipe(p) < 0)
    panic("pipe");
```

-----Point A-----

```
if(fork1() == 0){
```


```
    close(1);
    dup(p[1]);
    close(p[0]);
    close(p[1]);
```

-----Point B-----

```
    runcmd(pcmd->left);
```

```
}
```

fork() copies the descriptors too!



# pipe() and fork()

-----Point 0-----

case PIPE:

```
pcmd = (struct pipecmd*)cmd;
```

```
if(pipe(p) < 0)
```

```
    panic("pipe");
```

-----Point A-----

```
if(fork1() == 0){
```

```
    close(1);
```

```
    dup(p[1]);
```

Executed by child process

```
    close(p[0]);
```



```
    close(p[1]);
```

-----Point B-----

```
    runcmd(pcmd->left);
```

```
}
```

fork() copies the descriptors too!



# pipe() and fork()

-----Point 0-----

```
case PIPE:
pcmd = (struct pipecmd*)cmd;
if(pipe(p) < 0)
    panic("pipe");
```

-----Point A-----

```
if(fork1() == 0){
```

```
    close(1);
```

```
    dup(p[1]);
```

Executed by child process

```
    close(p[0]);
```



```
    close(p[1]);
```

-----Point B-----

```
    runcmd(pcmd->left);
```

```
}
```

fork() copies the descriptors too!



# pipe() and fork()

-----Point 0-----

```
case PIPE:
pcmd = (struct pipecmd*)cmd;
if(pipe(p) < 0)
    panic("pipe");
```

-----Point A-----



```
if(fork1() == 0){
    close(1);
    dup(p[1]); Executed by child process
    close(p[0]);
    close(p[1]);
```

-----Point B-----

```
    runcmd(pcmd->left);
}
```

fork() copies the descriptors too!

dup()'s destination is the lowest & unused file descriptor!



# pipe() and fork()

-----Point 0-----

```
case PIPE:
pcmd = (struct pipecmd*)cmd;
if(pipe(p) < 0)
    panic("pipe");
```



-----Point A-----

```
if(fork1() == 0){
    close(1);
    dup(p[1]);           Executed by child process
    close(p[0]);
    close(p[1]);
```

-----Point B-----

```
    runcmd(pcmd->left);
}
```

fork() copies the descriptors too!  
dup()'s destination is the lowest & unused file descriptor!



# pipe() and fork()

-----Point 0-----

```
case PIPE:
pcmd = (struct pipecmd*)cmd;
if(pipe(p) < 0)
    panic("pipe");
```

-----Point A-----



```
if(fork1() == 0){
    close(1);
    dup(p[1]);
    close(p[0]);
    close(p[1]);
```

-----Point B-----

```
    runcmd(pcmd->left);
}
```

Executed by child process

fork() copies the descriptors too!  
dup()'s destination is the lowest & unused file descriptor!





# pipe() and fork()

-----Point B-----

```
    runcmd(pcmd>left);
```

```
}
```

```
if(fork1() == 0){
```

```
    close(0);
```

```
    dup(p[0]);
```

```
    close(p[0]);
```

```
    close(p[1]);
```

```
    runcmd(pcmd->right);
```

```
}
```

```
close(p[0]);
```

```
close(p[1]);
```

-----Point C-----

```
wait();
```


```
wait();
```

```
break;
```

Executed by child process

fork() copies the descriptors too!

dup()'s destination is the lowest & unused file descriptor!

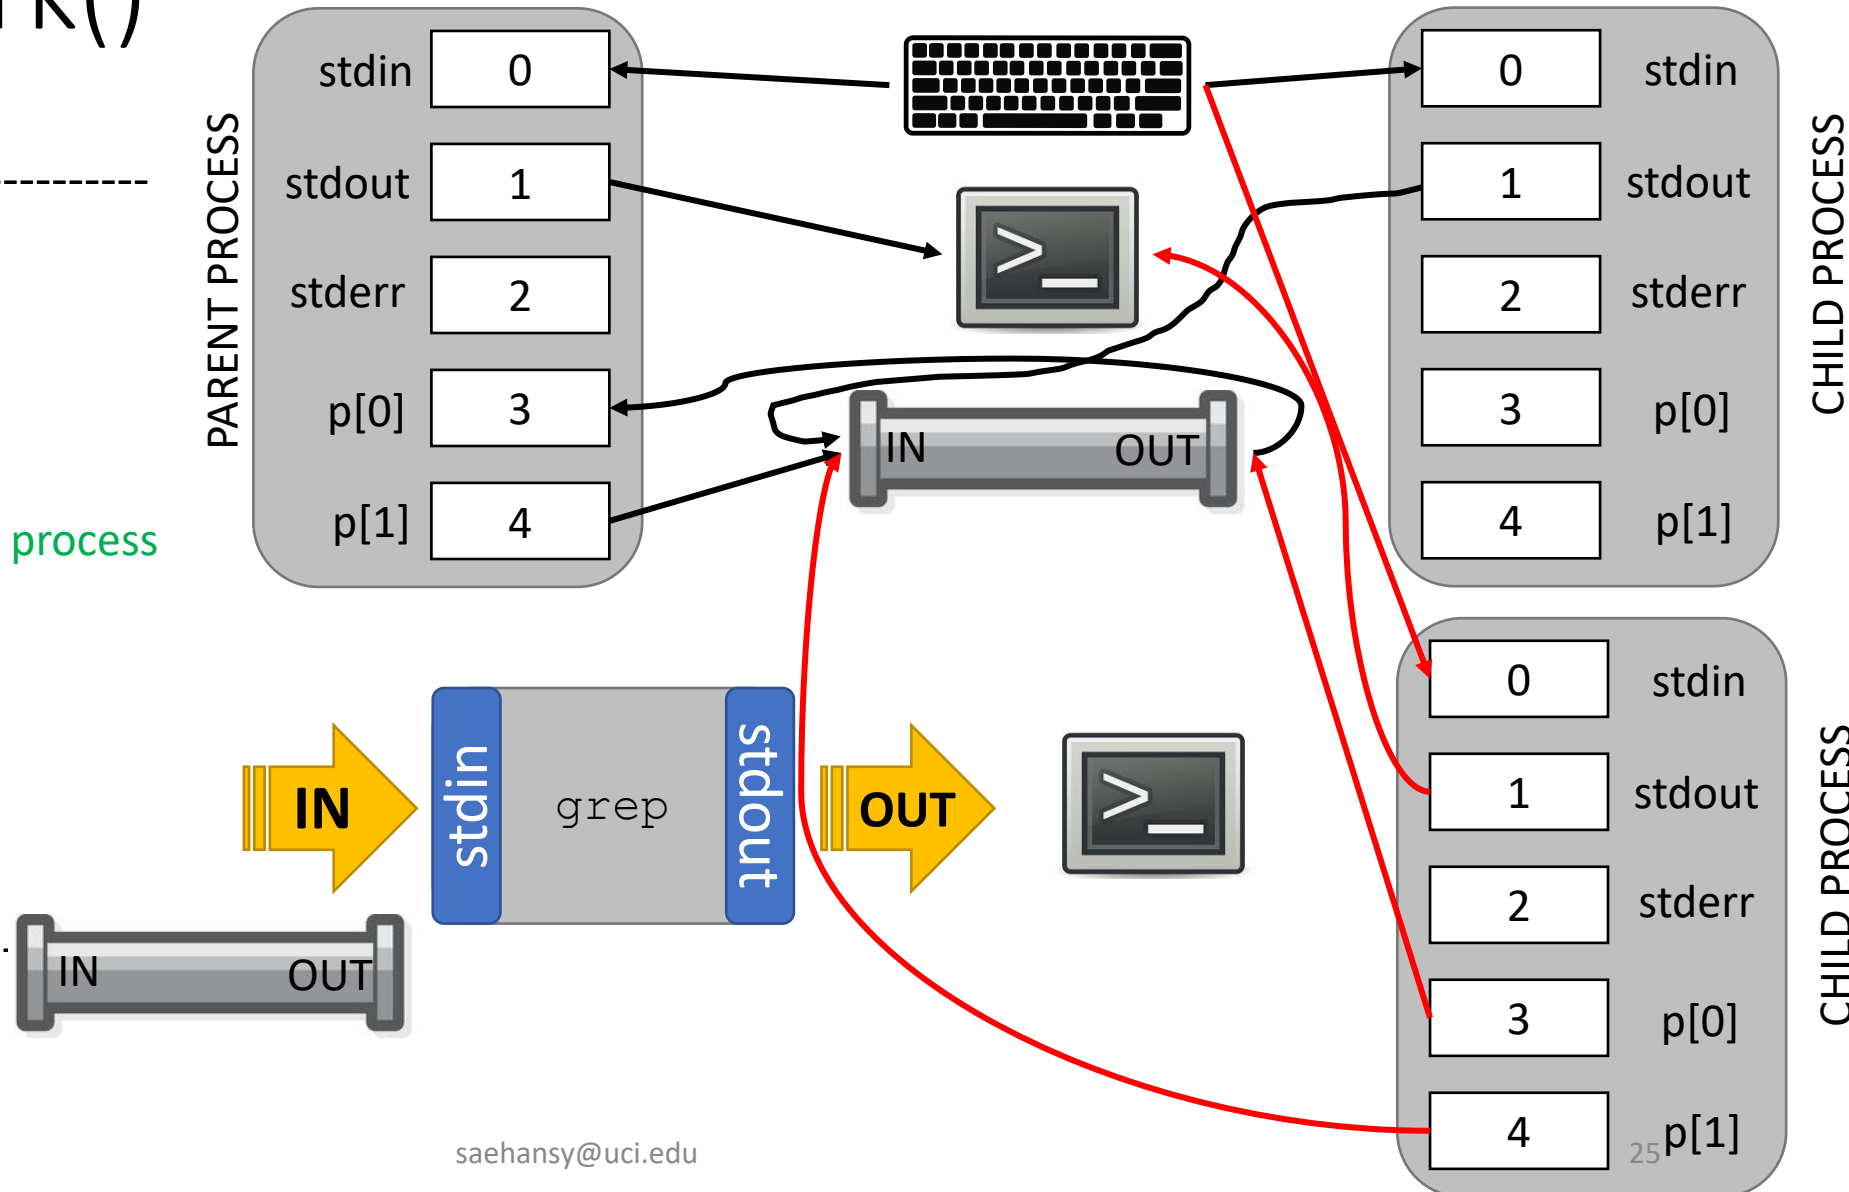


# pipe() and fork()

fork() copies the descriptors too!  
dup()'s destination is the lowest & unused file descriptor!

```

-----Point B-----
  runcmd(pcmd>left);
}
if(fork1() == 0){
  close(0);
  dup(p[0]);
  close(p[0]); Executed by child process
  close(p[1]);
  runcmd(pcmd->right);
}
close(p[0]);
close(p[1]);
-----Point C-----
wait();
wait();
break;
  
```

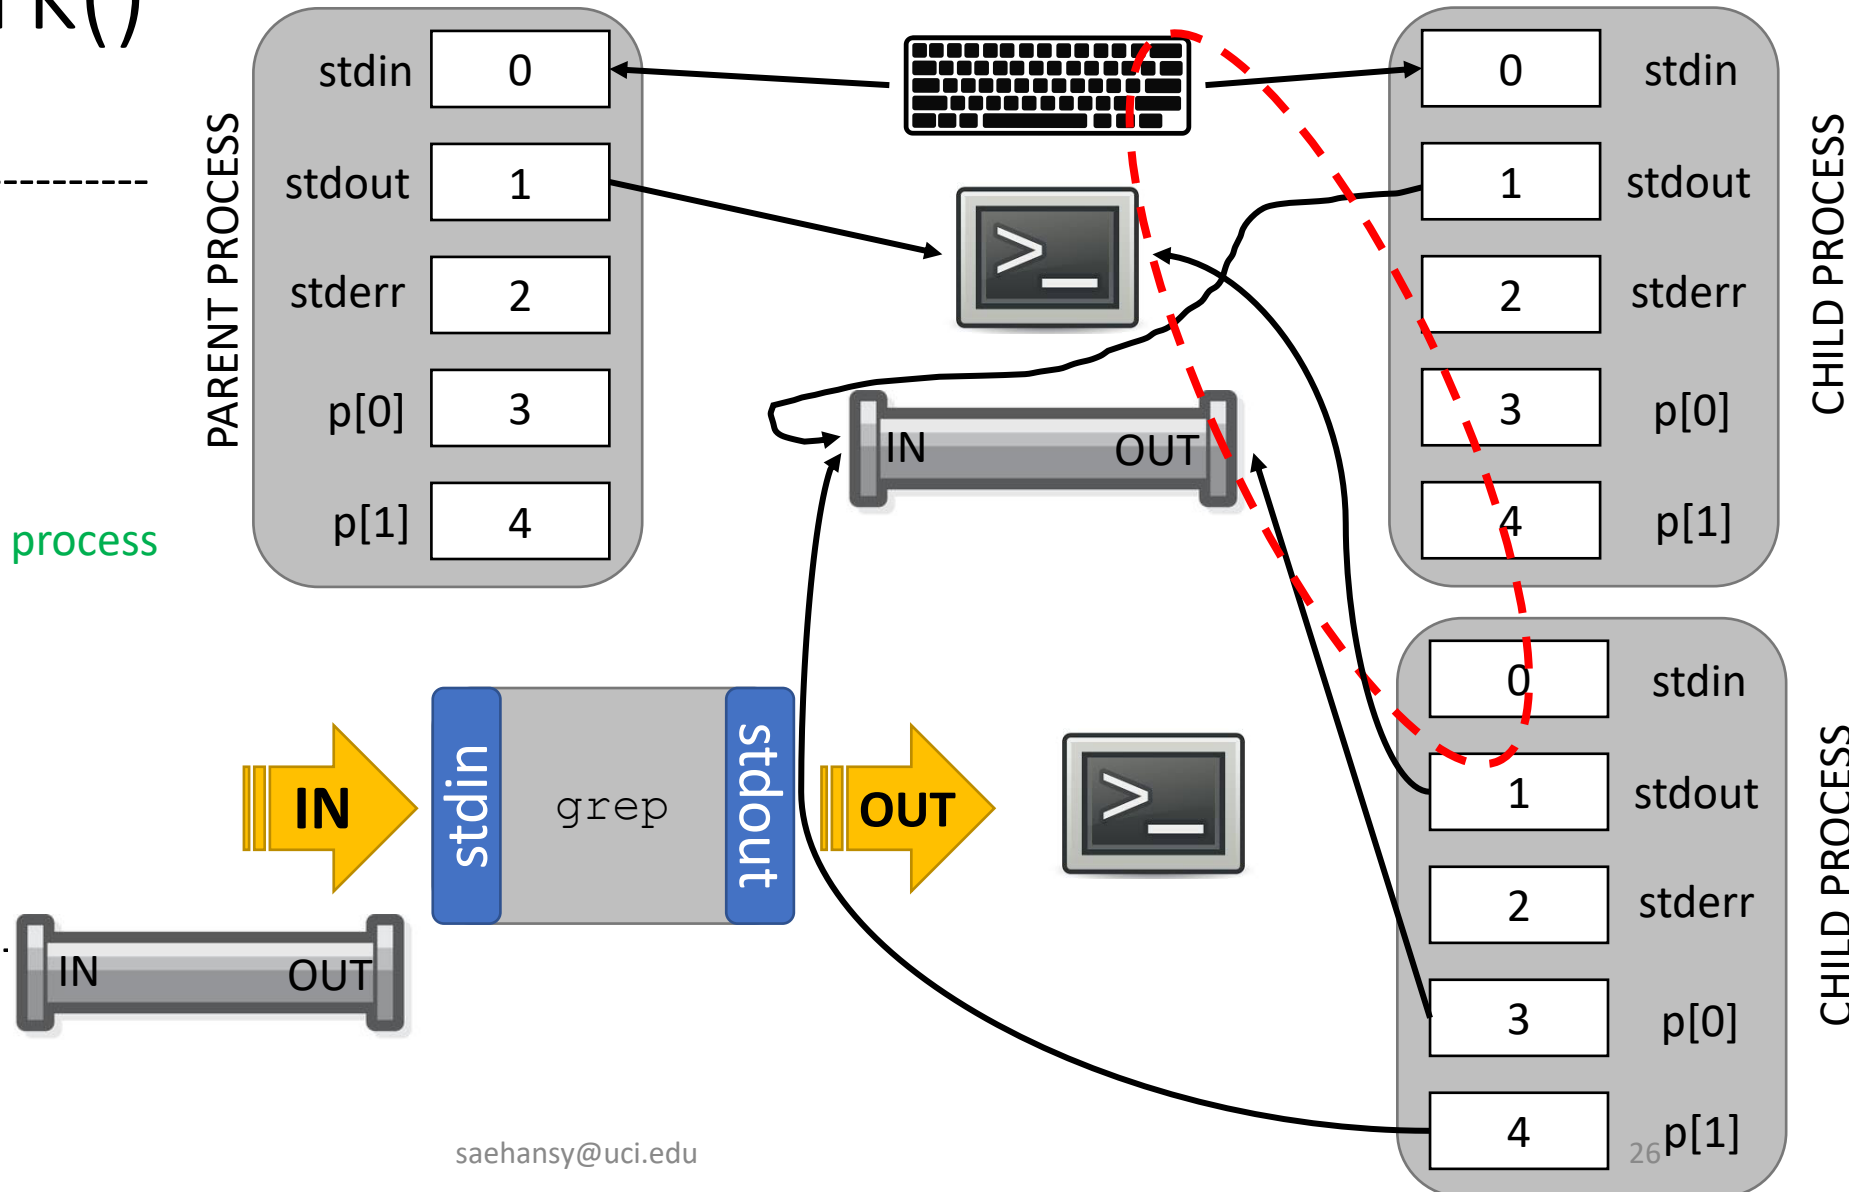


# pipe() and fork()

fork() copies the descriptors too!  
dup()'s destination is the lowest & unused file descriptor!

```
-----Point B-----  
    runcmd(pcmd>left);  
}  
if(fork1() == 0){  
    close(0);  
    dup(p[0]);  
    close(p[0]);  
    close(p[1]);  
    runcmd(pcmd->right);  
}  
close(p[0]);  
close(p[1]);  
-----Point C-----  
wait();  
wait();  
break;
```

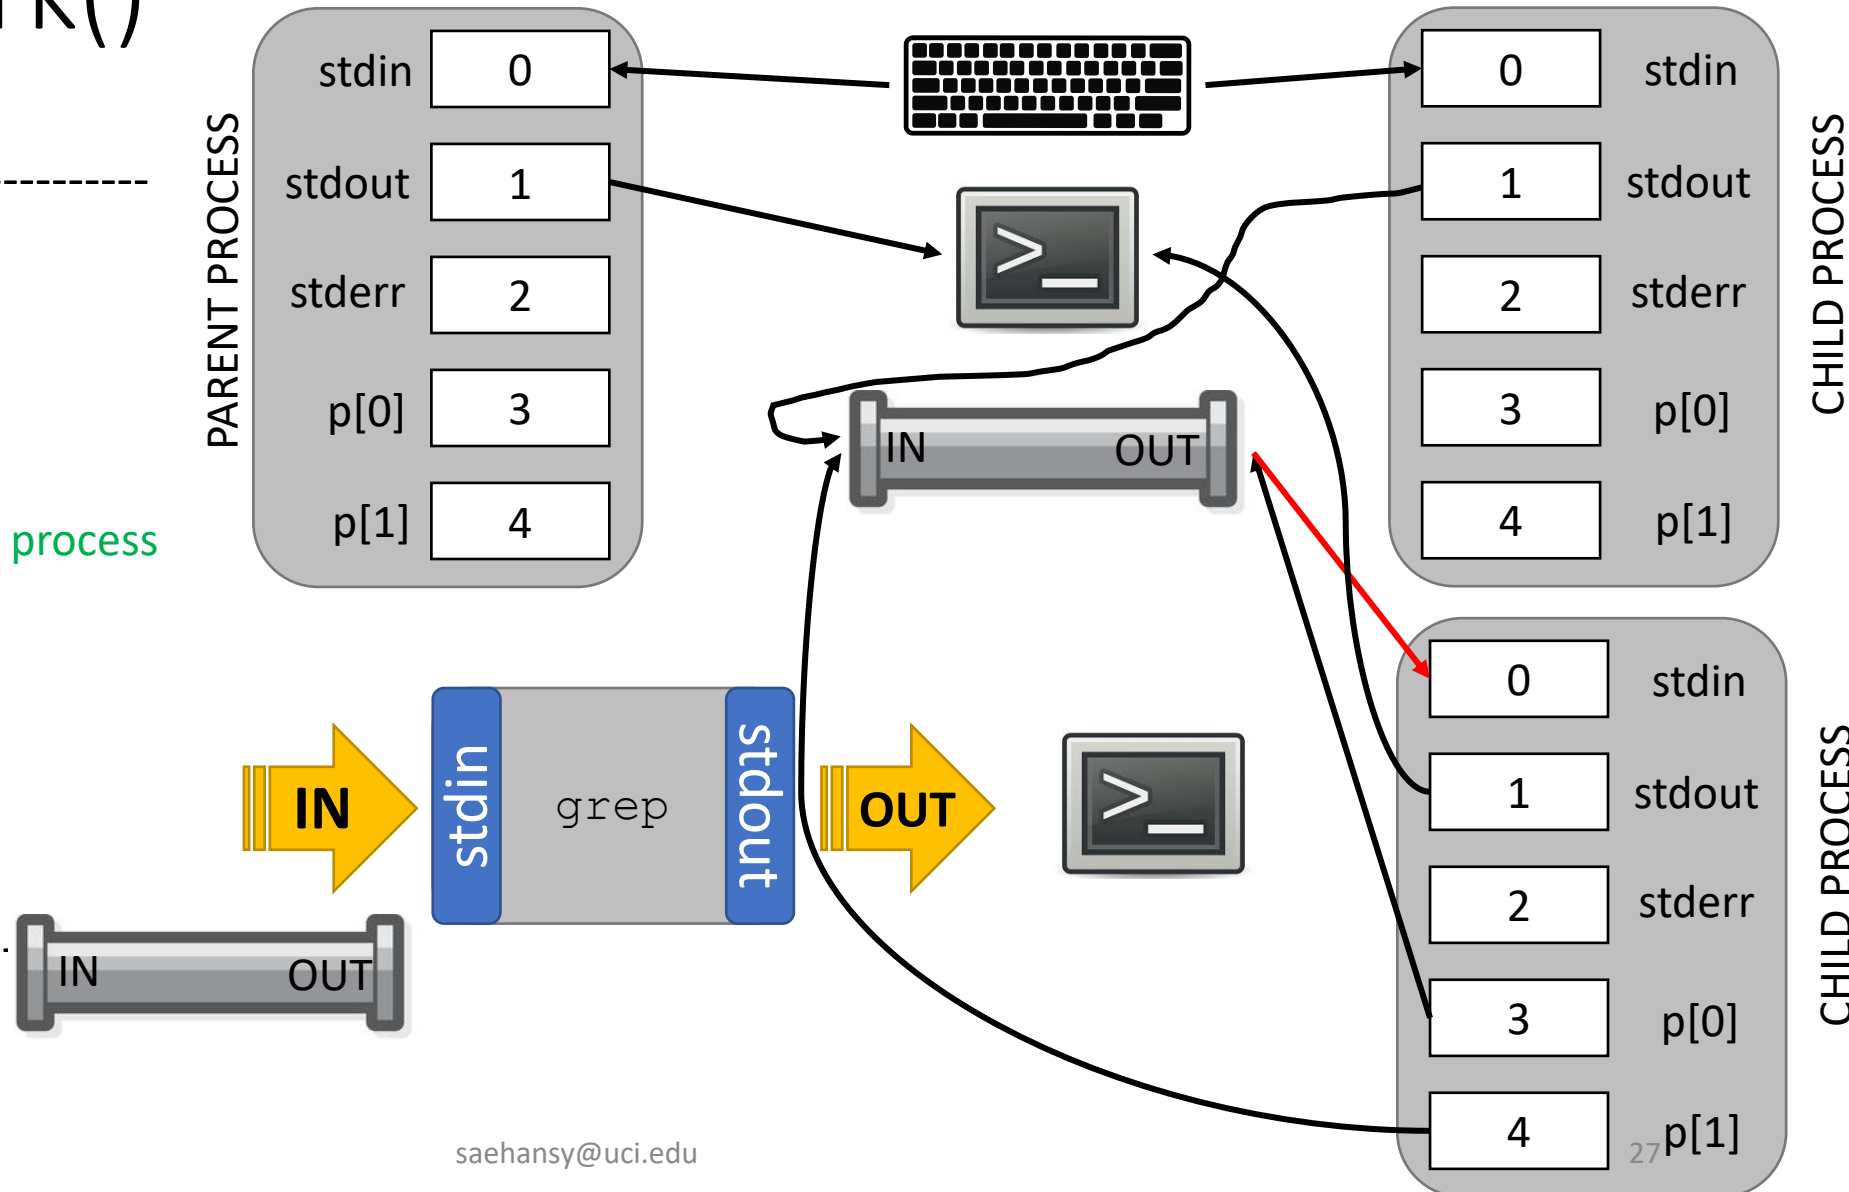
Executed by child process



# pipe() and fork()

fork() copies the descriptors too!  
dup()'s destination is the lowest & unused file descriptor!

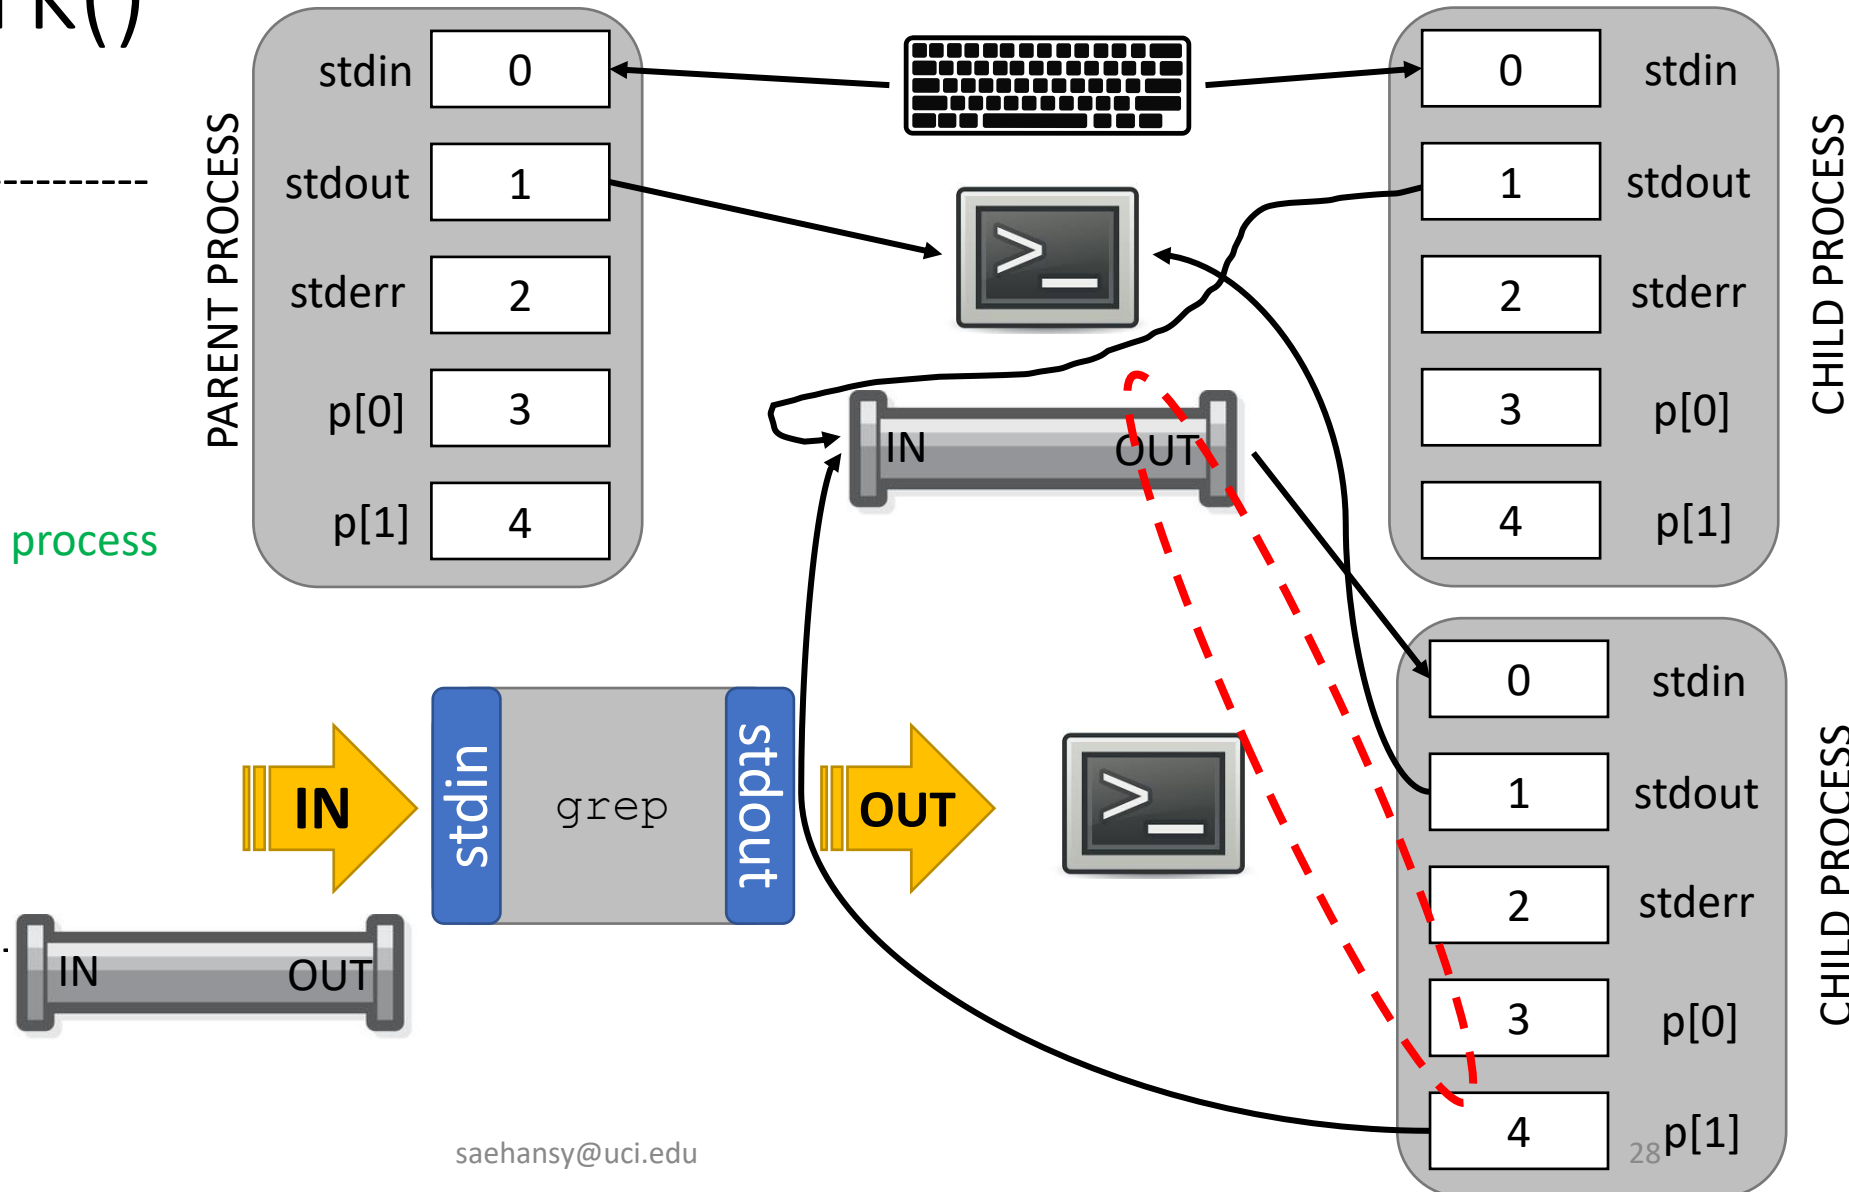
```
-----Point B-----  
    runcmd(pcmd>left);  
}  
if(fork1() == 0){  
    close(0);  
    dup(p[0]);  
    close(p[0]); Executed by child process  
    close(p[1]);  
    runcmd(pcmd->right);  
}  
close(p[0]);  
close(p[1]);  
-----Point C-----  
wait();  
wait();  
break;
```



# pipe() and fork()

fork() copies the descriptors too!  
dup()'s destination is the lowest & unused file descriptor!

```
-----Point B-----  
    runcmd(pcmd>left);  
}  
if(fork1() == 0){  
    close(0);  
    dup(p[0]);  
    close(p[0]); Executed by child process  
    close(p[1]);  
    runcmd(pcmd->right);  
}  
close(p[0]);  
close(p[1]);  
-----Point C-----  
wait();  
wait();  
break;
```

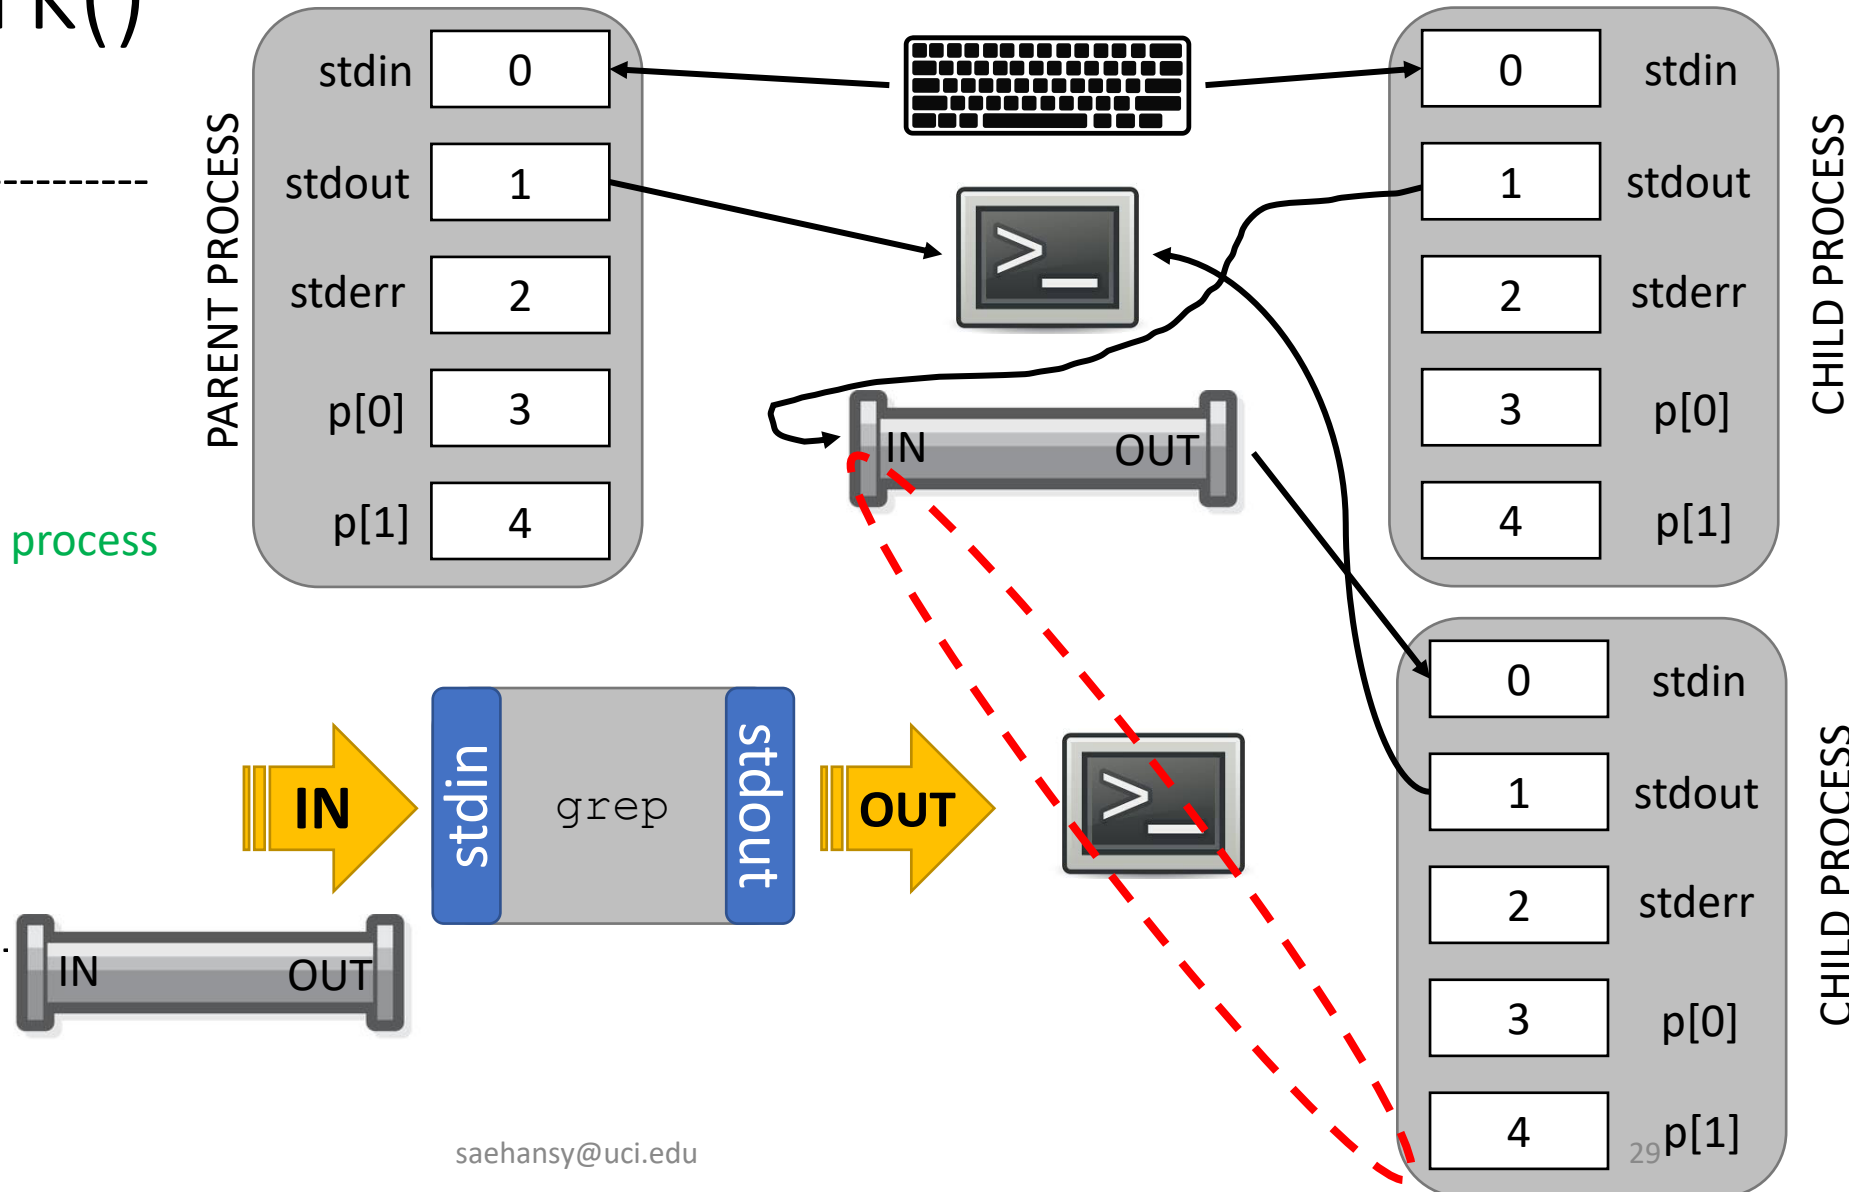


# pipe() and fork()

fork() copies the descriptors too!  
dup()'s destination is the lowest & unused file descriptor!

```
-----Point B-----  
    runcmd(pcmd>left);  
}  
if(fork1() == 0){  
    close(0);  
    dup(p[0]);  
    close(p[0]);  
    close(p[1]);  
    runcmd(pcmd->right);  
}  
close(p[0]);  
close(p[1]);  
-----Point C-----  
wait();  
wait();  
break;
```

Executed by child process



# pipe() and fork()

-----Point B-----

```
    runcmd(pcmd>left);
```

```
}
```

```
if(fork1() == 0){
```

```
    close(0);
```

```
    dup(p[0]);
```

```
    close(p[0]);
```

```
    close(p[1]);
```

```
    runcmd(pcmd->right);
```

```
}
```

```
close(p[0]);
```

```
close(p[1]);
```

-----Point C-----

```
wait();
```


```
wait();
```

```
break;
```

Parent waits child processes

fork() copies the descriptors too!

dup()'s destination is the lowest & unused file descriptor!




# pipe() and fork() and exec()

```
if(fork1() == 0){  
    ...  
    runcmd(pcmd->right);  
}
```

*runcmd() contains exec functions*

```
$  
$ ls | grep asdf  
asdfasdf  
$
```



**`int execvp(const char *file, char *const argv[]);`**  
replaces the current process image with a new process image.



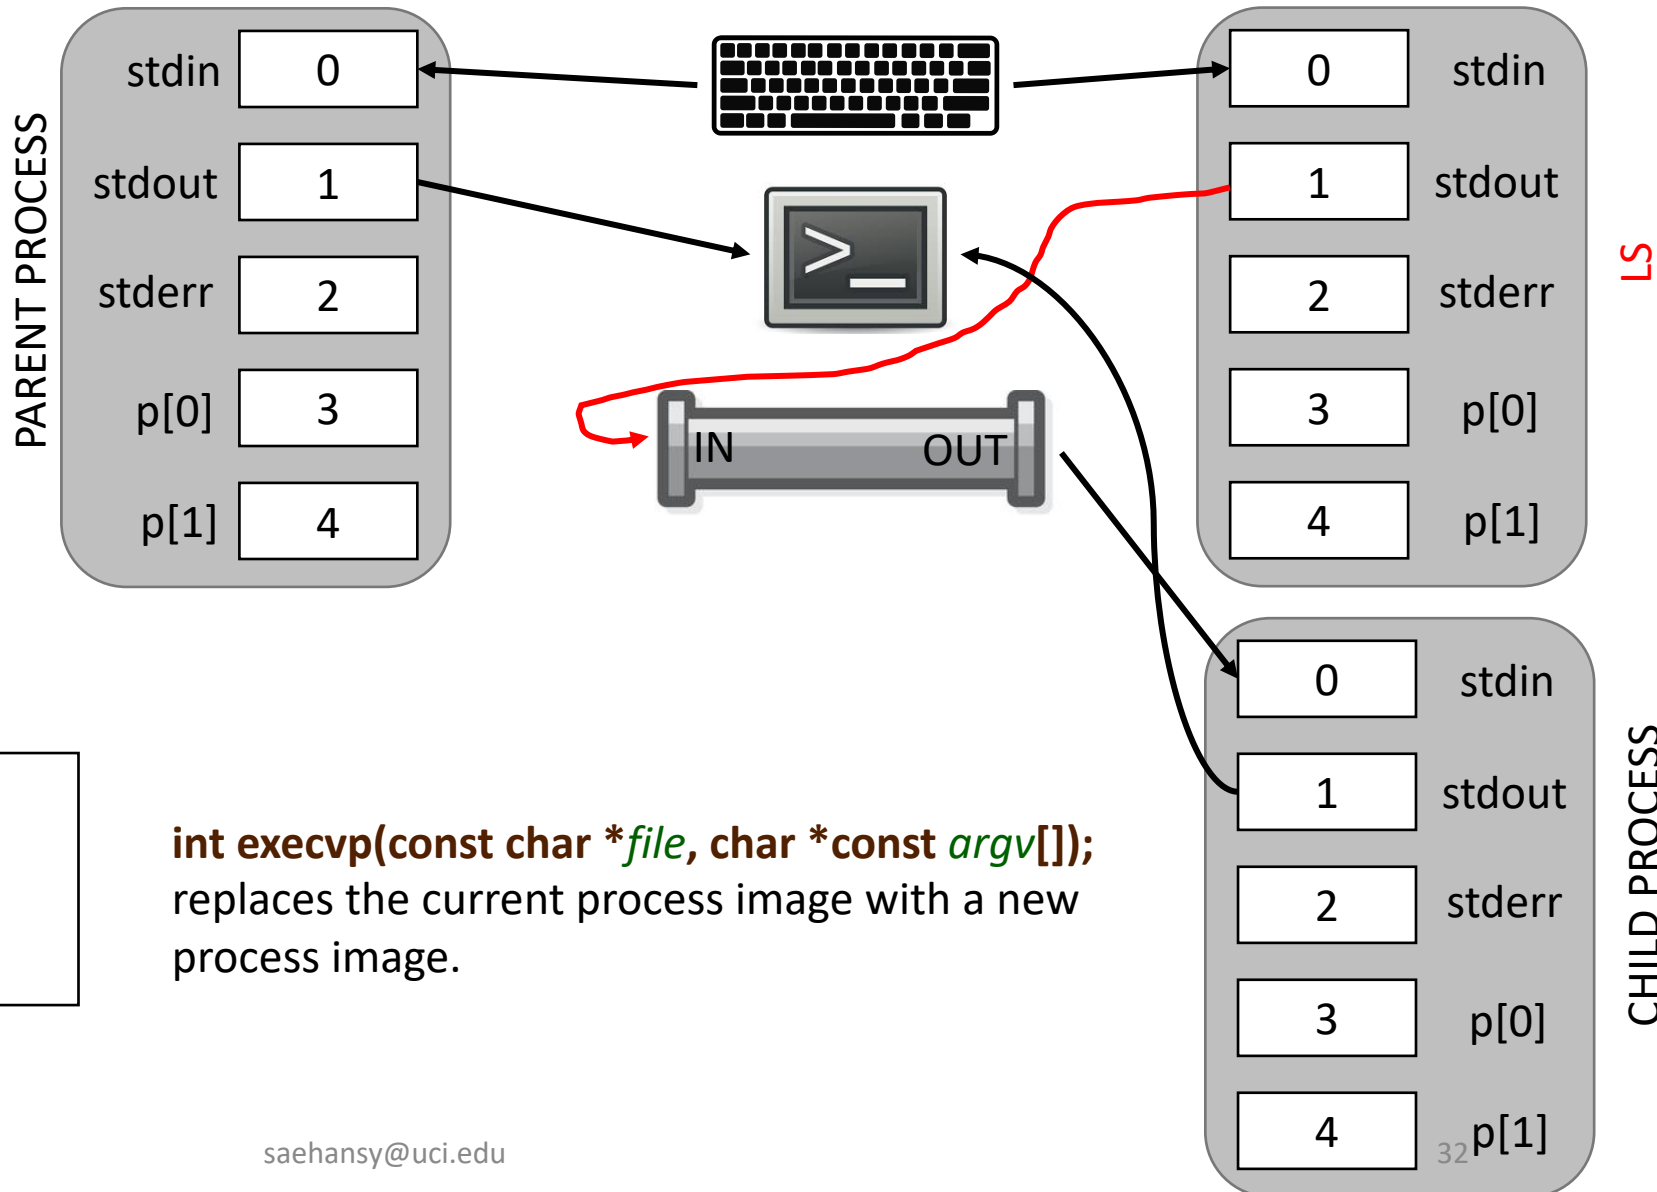
# pipe() and fork() and exec()

```
if(fork1() == 0){
    ...
    runcmd(pcmd->right);
}
```

*runcmd() contains exec functions*

```
$  
$ ls | grep asdf  
asdfasdf  
$
```

**int execvp(const char \**file*, char \*const *argv*[]);**  
replaces the current process image with a new process image.

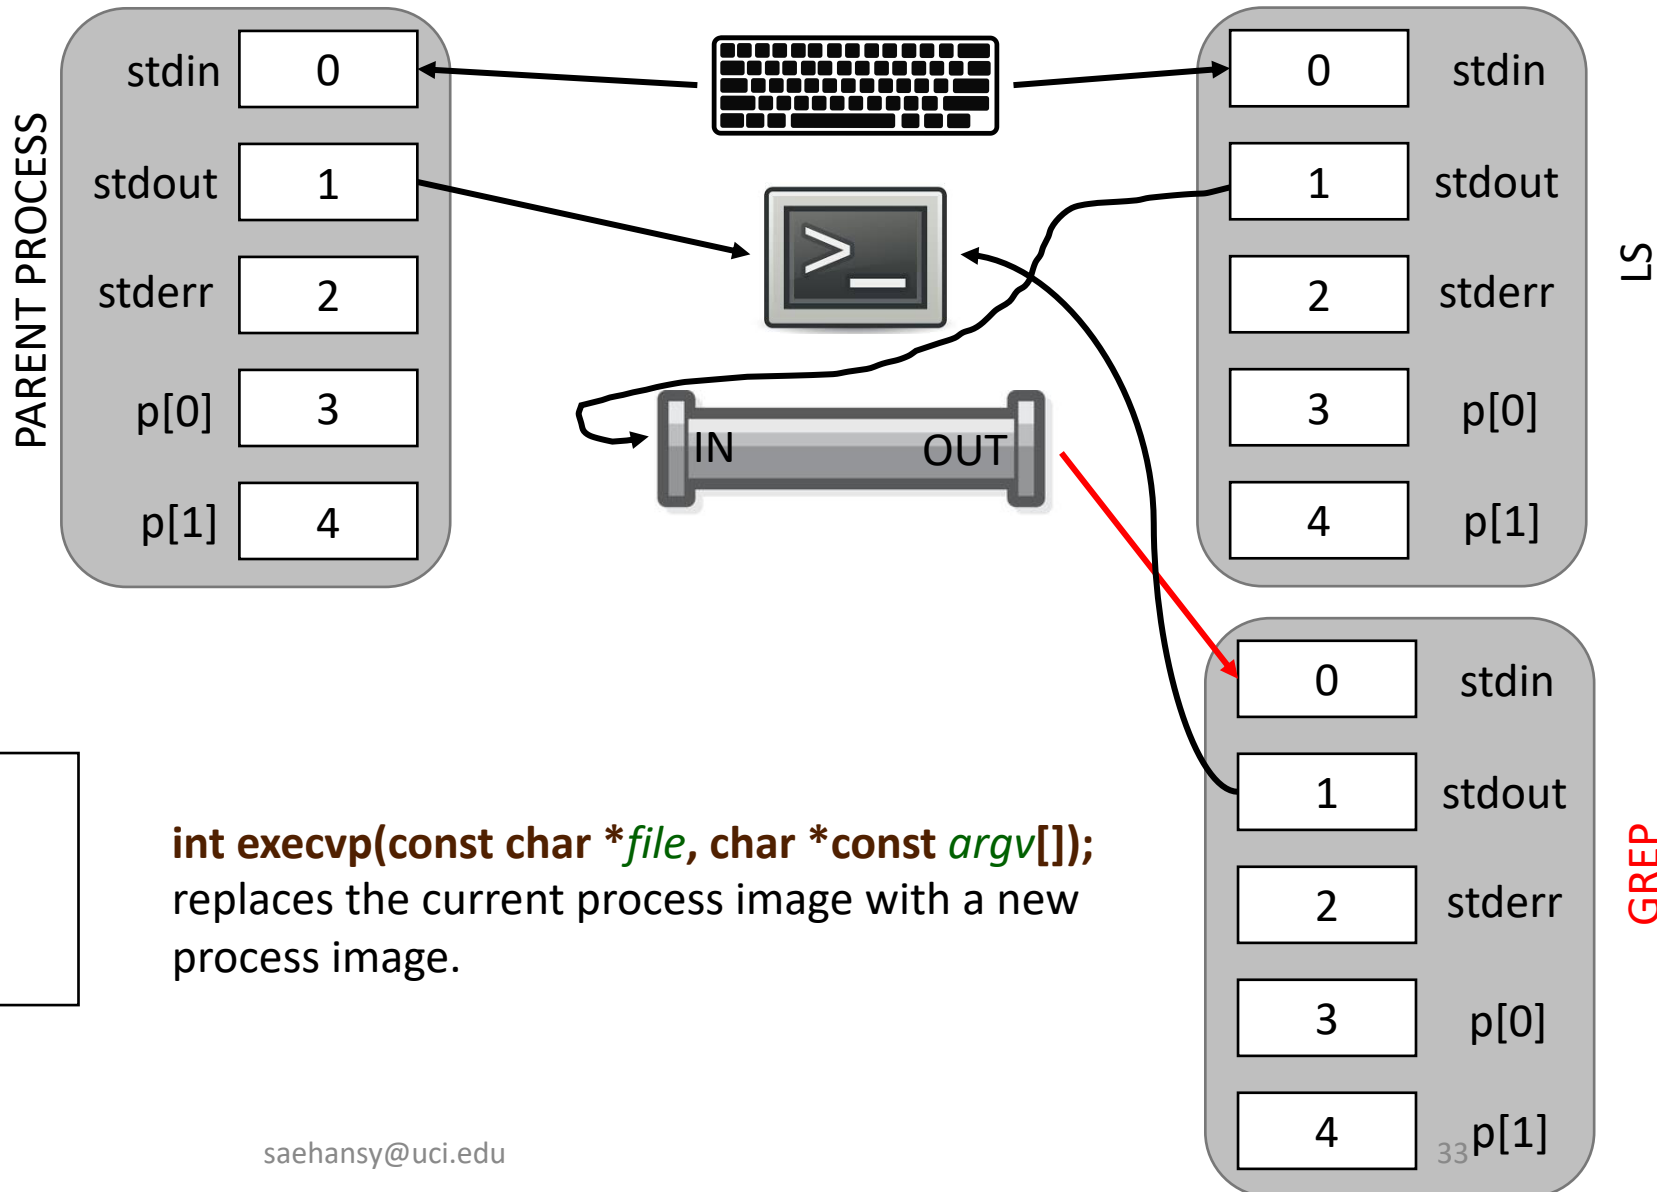


# pipe() and fork() and exec()

```
if(fork1() == 0){  
    ...  
    runcmd(pcmd->right);  
}
```

*runcmd() contains exec functions*

```
$  
$ ls | grep asdf  
asdfasdf  
$
```

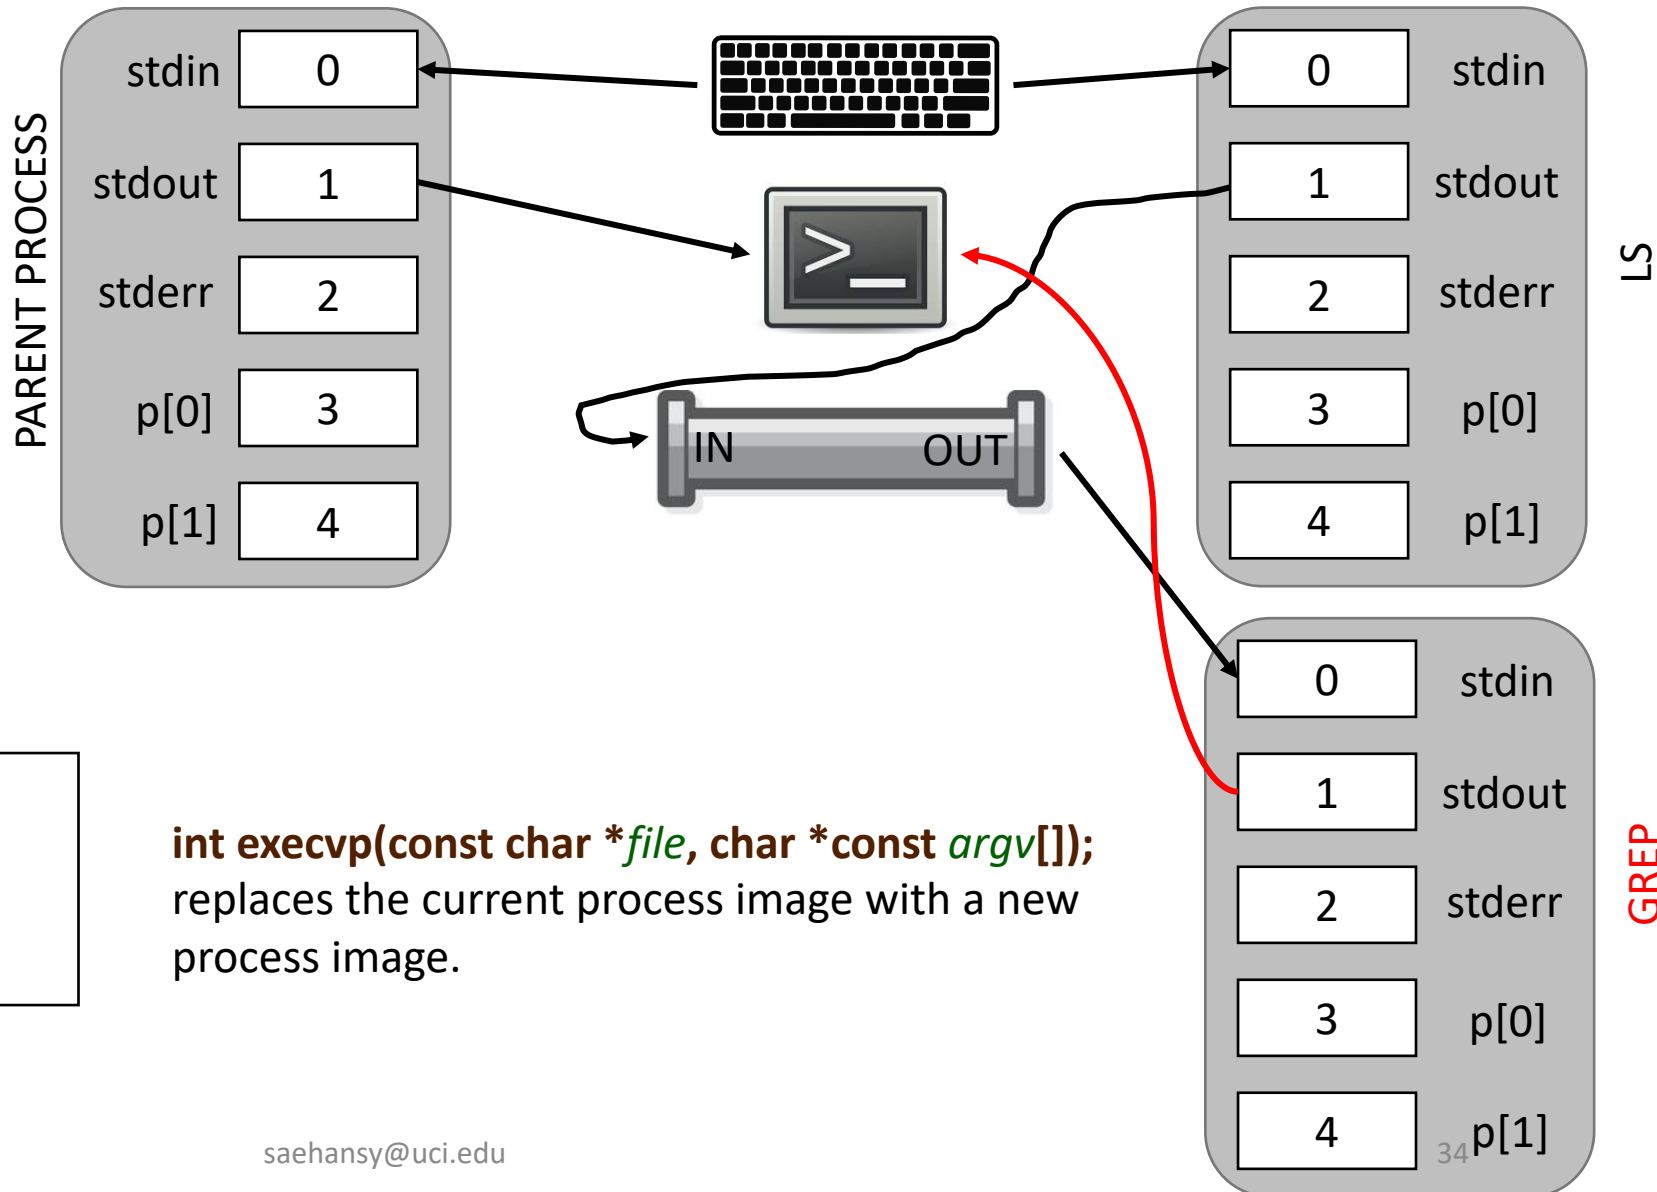


# pipe() and fork() and exec()

```
if(fork1() == 0){  
    ...  
    runcmd(pcmd->left);  
}
```

*runcmd() contains exec functions*

```
$  
$ ls | grep asdf  
asdfasdf  
$
```



**int execvp(const char \*file, char \*const argv[]);**  
replaces the current process image with a new process image.