

Principles of Operating Systems
Fall 2017
Final
12/13/2017
Time Limit: 8:00am - 10:00am

Name (Print): _____

- Don't forget to write your name on this exam.
- This is an open book, open notes exam. But no online or in-class chatting.
- Ask me if something is not clear in the questions.
- **Organize your work**, in a reasonably neat and coherent way, in the space provided. Work scattered all over the page without a clear ordering will receive very little credit.
- **Mysterious or unsupported answers will not receive full credit.** A correct answer, unsupported by explanation will receive no credit; an incorrect answer supported by substantially correct explanations might still receive partial credit.
- If you need more space, use the back of the pages; clearly indicate when you have done this.

Problem	Points	Score
1	20	
2	5	
3	10	
4	15	
5	20	
6	15	
7	5	
8	5	
Total:	95	

1. File system

Xv6 lays out the file system on disk as follows:

super	log header	log	inode	bmap	data
1	2	3	32	58	59

Block 1 contains the super block. Blocks 2 through 31 contain the log header and the log. Blocks 32 through 57 contain inodes. Block 58 contains the bitmap of free blocks. Blocks 59 through the end of the disk contain data blocks.

Ben modifies the function `bwrite` in `bio.c` to print the block number of each block written.

Ben boots xv6 with a fresh `fs.img` and types in the command `rm README`, which deletes the `README` file. This command produces the following trace:

```
$ rm README
write 3
write 4
write 5
write 2
write 59
write 32
write 58
write 2
$
```

(a) (5 points) Briefly explain what block 59 contains in the above trace. Why is it written?

(b) (5 points) What does block 5 contain? Why is it written?

- (c) (10 points) How many non-zero bytes are written to block 2 when it's written the first time and what are the bytes? (To get the full credit you have to explain what block 2 contains, and why each non-zero byte is written).

2. Synchronization

- (a) (5 points) Ben runs xv6 on a single CPU machine, he decides it's a good idea to get rid of the `acquire()` and `release()` functions, since after all they take some time but seem unnecessary in a single-CPU scenario. Explain if removal of these functions is fine.

3. Process memory layout

Bob decides to implement the following xv6 program (hello)

```
#define PGSIZE 4096
int main(int argc, char *argv[]) {
    char buf[PGSIZE] = {0};
    printf(1, "Hello World!, %p\n", buf);
    exit();
}
```

(a) (10 points) When Bob runs it he encounters the following error message:

```
pid 3 hello: trap 14 err 7 on cpu 1 eip 0x22 addr 0x1fd0--kill proc
```

Can you help Bob understand the problem? (To receive full points, you should explain why this error happened, what does the error values mean and how Bob can fix his code)

4. Virtual memory

Ben wants to know the address of physical pages that back up virtual memory of his process. He digs into the kernel source and comes across the `V2P()` macro that is frequently used in the kernel.

```
#define KERNBASE 0x80000000
#define V2P(a) (((uint) (a)) - KERNBASE)
```

He decides to try the `V2P` macro in his program (below), but encounters a crash.

```
void test(void) {
    int a;
    *(uint*)V2P(&a) = 0xaddb;
    printf(1, "I changed physical memory at %x\n", a);
}
int main(int argc, char *argv[]) {
    test();
    exit();
}
```

(a) (5 points) Explain what is going on and why Ben's program crashes.

(b) (10 points) Ben puts the code of the `test()` function inside a new system call trying to see if it works inside the kernel. Will it work (explain your answer)?

5. Demand paging

Ben wants to extend xv6 with demand paging. Ben observes that some pages of user processes (heap, text, and stack) are not accessed that frequently, yet anyway they consume valuable physical memory. So Ben comes up with a plan to free these infrequently used pages by saving them to disk (swapping). For an idle page he plans to unmap it from the process page table, save content of the page to disk (i.e., in a special swap area on disk), and free the physical page back to the kernel, making it accessible for other processes. Obviously, Ben wants paging to be transparent. I.e., when a process accesses one of the swapped pages, Ben plans to catch an exception, allocate a new physical page, read old content of the page from disk, and fix the process page table in such a way that process can access the page like nothing happened.

- (a) (5 points) How can Ben unmap a page from the process address space? I.e., what changes to the process page table are required to catch an exception when the process tries to access a swapped page (hint: look at how guard page is implemented)?

- (b) (15 points) Ben plans to catch the exception caused by an unmapped page access inside the `trap()` function. Provide a sketch for the code that implements the exception handling, reads page from disk, and maps it back into the process address space.

6. Process creation

While editing the xv6 code, Jimmy accidentally erases the below section of code under `fork()` function on `proc.c`

```
2584 for(i = 0; i < NOFILE; i++)
2585     if(proc->ofile[i])
2586         np->ofile[i] = filedup(proc->ofile[i]);
```

(a) (5 points) Explain what the above section of code does?

(b) (10 points) Explain where things can go wrong without this code. Quote a concrete example.

7. System call API

Alice executes the following program

```
main() {  
    char *msg = "foo\n";  
    int pid = fork();  
    if (pid){  
        msg = "bar\n";  
    } else {  
        msg = "baz\n";  
    }  
    write(1, msg, 4);  
    exit(0);  
}
```

- (a) (5 points) What are all possible outputs of this program? Explain your answer.

8. ics143A. I would like to hear your opinions about 6.828, so please answer the following questions.
(Any answer, except no answer, will receive full credit.)

(a) (1 point) Grade ics143A on a scale of 0 (worst) to 10 (best)?

(b) (2 points) Any suggestions for how to improve ics143A?

(c) (1 point) What is the best aspect of ics143A?

(d) (1 point) What is the worst aspect of ics143A?