

Principles of Operating Systems
Fall 2017
Final
12/13/2017
Time Limit: 8:00am - 10:00am

Name (Print): _____

- Don't forget to write your name on this exam.
- This is an open book, open notes exam. But no online or in-class chatting.
- Ask me if something is not clear in the questions.
- **Organize your work**, in a reasonably neat and coherent way, in the space provided. Work scattered all over the page without a clear ordering will receive very little credit.
- **Mysterious or unsupported answers will not receive full credit.** A correct answer, unsupported by explanation will receive no credit; an incorrect answer supported by substantially correct explanations might still receive partial credit.
- If you need more space, use the back of the pages; clearly indicate when you have done this.

Problem	Points	Score
1	20	
2	10	
3	15	
4	20	
5	15	
6	5	
Total:	85	

1. File system

Xv6 lays out the file system on disk as follows:

super	log header	log	inode	bmap	data
1	2	3	32	58	59

Block 1 contains the super block. Blocks 2 through 31 contain the log header and the log. Blocks 32 through 57 contain inodes. Block 58 contains the bitmap of free blocks. Blocks 59 through the end of the disk contain data blocks.

Ben modifies the function `bwrite` in `bio.c` to print the block number of each block written.

Ben boots xv6 with a fresh `fs.img` and types in the command `rm README`, which deletes the `README` file. This command produces the following trace:

```
$ rm README
write 3
write 4
write 5
write 2
write 59
write 32
write 58
write 2
$
```

- (a) (5 points) Briefly explain what block 59 contains in the above trace. Why is it written?

(b) (5 points) What does block 5 contain? Why is it written?

(c) (10 points) How many non-zero bytes are written to block 2 when it's written the first time and what are the bytes? (To get the full credit you have to explain what block 2 contains, and why each non-zero byte is written).

2. Memory management.

(a) (5 points) Explain organization of the xv6 memory allocator.

(b) (5 points) Why do you think xv6 does not have buddy or slab allocators? Under what conditions you would have to add these allocators to the xv6 kernel?

3. Synchronization

- (a) (5 points) Ben runs xv6 on a single CPU machine, he decides it's a good idea to get rid of the `acquire()` and `release()` functions, since after all they take some time but seem unnecessary in a single-CPU scenario. Explain if removal of these functions is fine.

- (b) (10 points) Alyssa runs xv6 on a machine with 8 processors and 8 processes. Each process calls `uptime()` (3738) system call continuously, reading the number of ticks passed since boot. Alyssa measures the number of `uptime()` system calls per second and notices that 8 processes achieve the same total throughput as 1 process, even though each process runs on a different processor. Why is the throughput of 8 processes the same as that of 1 process?

4. Scheduling

- (a) (10 points) You would like to extend xv6 with priority based scheduler, i.e., each process has a priority, and processes with a higher priority are scheduled first. Write the code for your implementation below (which xv6 functions need to be changed?)

- (b) (10 points) Now you would like to extend your priority scheduler with support for interactive tasks, e.g., a task that spends a lot of time waiting, should run first (i.e., receive priority boost). Provide code that handles waiting tasks and implements priority boost (again, just change related xv6 functions).

5. Page tables.

Xv6 uses 4MB page table during boot. It is defined as:

```
1406 // The boot page table used in entry.S and entryother.S.
1407 // Page directories (and page tables) must start on page boundaries,
1408 // hence the __aligned__ attribute.
1409 // PTE_PS in a page directory entry enables 4Mbyte pages.
1410
1411 __attribute__((__aligned__(PGSIZE)))
1412 pde_t entrypgdir[NPDENTRIES] = {
1413     // Map VAs [0, 4MB) to PAs [0, 4MB)
1414     [0] = (0) | PTE_P | PTE_W | PTE_PS,
1415     // Map VAs [KERNBASE, KERNBASE+4MB) to PAs [0, 4MB)
1416     [KERNBASE >> PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
1417 };
```

- (a) (5 points) What virtual addresses (and to what physical addresses) does this page table map?

- (b) (10 points) Imagine now that 4MB pages are not available, and you have to use regular 4KB pages. How do you need to change the definition of entrypgdir for xv6 to work correctly (provide code and short explanation).

6. ics143A. I would like to hear your opinions about 6.828, so please answer the following questions.
(Any answer, except no answer, will receive full credit.)

(a) (1 point) Grade ics143A on a scale of 0 (worst) to 10 (best)?

(b) (2 points) Any suggestions for how to improve ics143A?

(c) (1 point) What is the best aspect of ics143A?

(d) (1 point) What is the worst aspect of ics143A?