

COMPSCI: Principles of Operating Systems

Lecture 1: Introduction

Anton Burtsev
January, 2017

Class details

- Undergraduate
 - 22 students
- Instructor: Anton Burtsev
- Meeting time: 9:00-9:50am (M, W, F)
 - Discussions: 1:00-1:50am (F)
 - No regular discussion section unless scheduled
 - Feel free to stop by my office with questions (DBH 3066)
- No TAs
- Web page
 - <http://www.ics.uci.edu/~aburtsev/143A>

This course

- Inspired by
 - MIT 6.828: Operating System Engineering
 - <https://pdos.csail.mit.edu/6.828/2016/>
 - Adapted for undergraduate students
- We will use xv6
 - Relatively simple (9K lines of code)
 - Reasonably complete UNIX kernel
 - <https://pdos.csail.mit.edu/6.828/2016/xv6.html>
- xv6 comes with a book
 - <https://pdos.csail.mit.edu/6.828/2016/xv6/book-rev9.pdf>
- And source code printout
 - <https://pdos.csail.mit.edu/6.828/2016/xv6/xv6-rev9.pdf>

Course Book

“Operating Systems: Three Easy Pieces”
Remzi H. Arpaci-Dusseau and Andrea C.
Arpaci-Dusseau

- Free online version

<http://pages.cs.wisc.edu/~remzi/OSTEP/>

Course organization


- Lectures
 - High level concepts and abstractions
- Reading
 - Xv6 book + source code
 - OSTEP book
- Homeworks
 - Coding real parts of the xv6 kernel
- Design riddles
 - Understanding design tradeoffs, explaining parts of xv6

Prerequisites


- Solid C coding skills
 - Xv6 is written in C
 - You need to read, code and debug
 - All homeworks are in C
 - Many questions will require explaining xv6 code
- Be able to work and code in Linux/UNIX
- Some assembly skills

What is an operating system?

Goal: Run your code on a piece of hardware



- Read CPU manual
- A tiny boot layer
 - Initialize CPU, memory
 - Jump to your code
- `main()`
 - **This is your OS!**




Print out a string

- On the screen or serial line


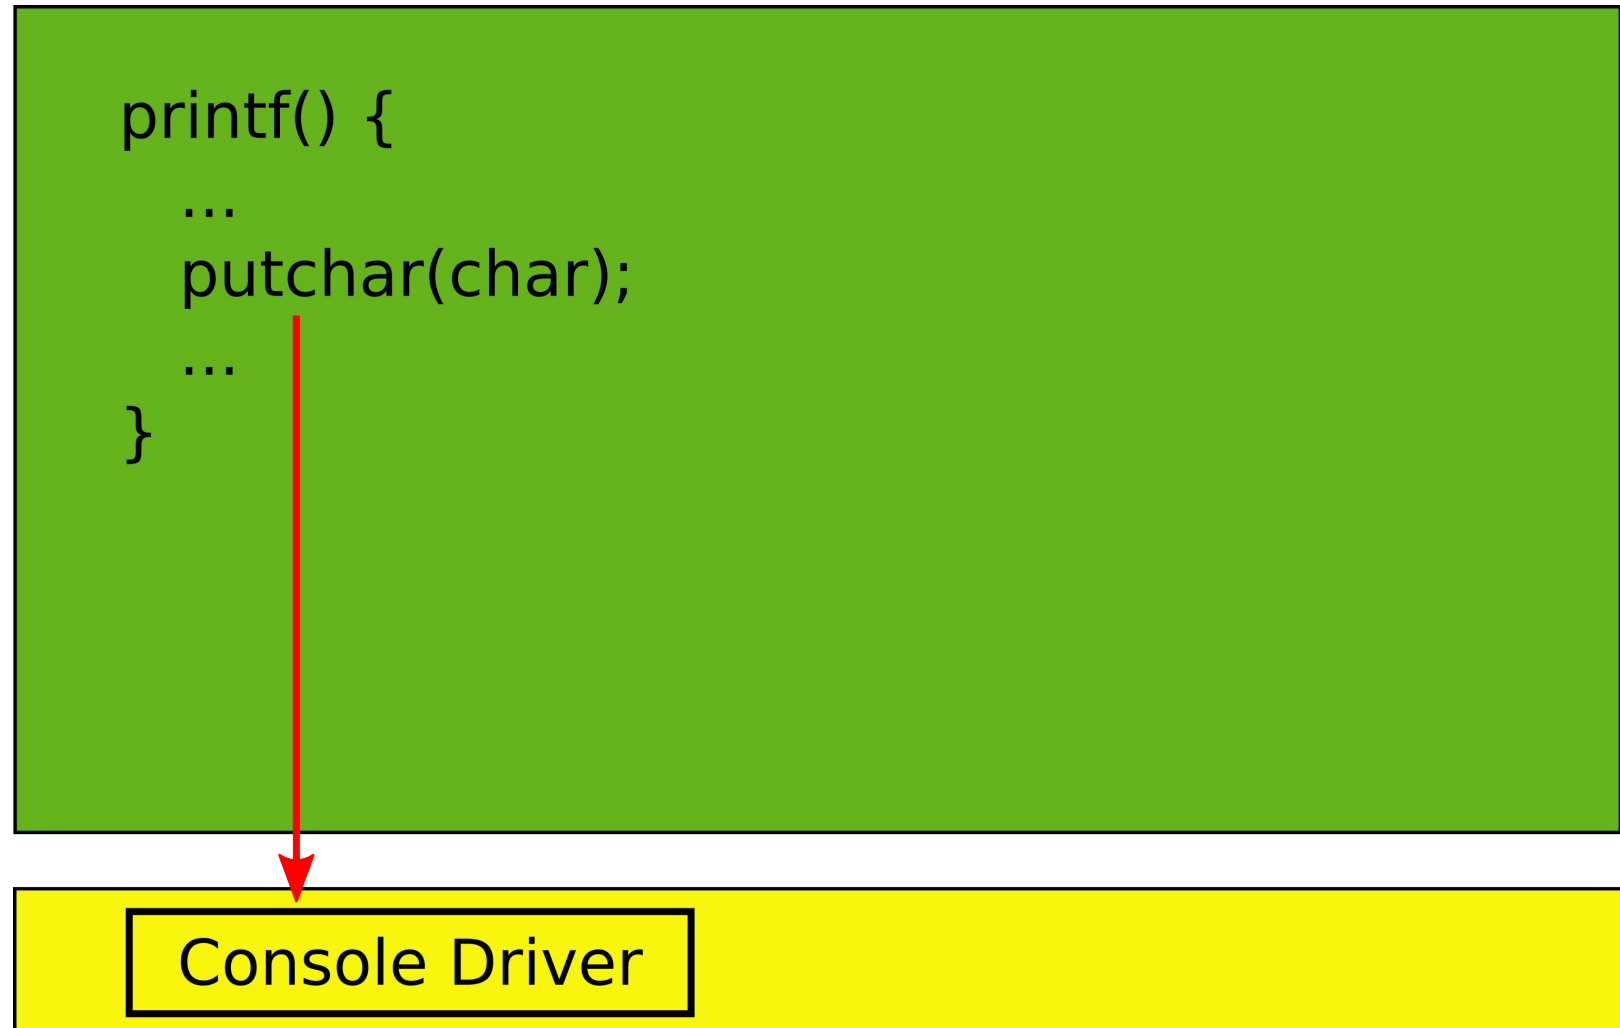
```
printf() {  
    ...  
    if (vga) {  
        asm("mov <magic number 1>, char");  
    } else if (serial) {  
        asm("out <magic number 2>, char");  
    }  
    ...  
}
```

OS



A more general interface

- First device driver



Device drivers

- Abstract hardware
 - Provide high-level interface
 - Hide minor differences
 - Implement some optimizations
 - Batch requests
- Examples
 - Console, disk, network interface
 - ...virtually any piece of hardware you know


Goal: Want to run two programs

```
main() {  
    ...  
    yield()  
}
```


```
main() {  
    ...  
    yield()  
}
```

- What does it mean?
 - Only one CPU
- Run one, then run another one

Save/restore




Goal: Want to run two programs



- Exit into the kernel periodically
- Context switch
 - Save and restore context
 - Essentially registers

- What! Two programs, one memory?

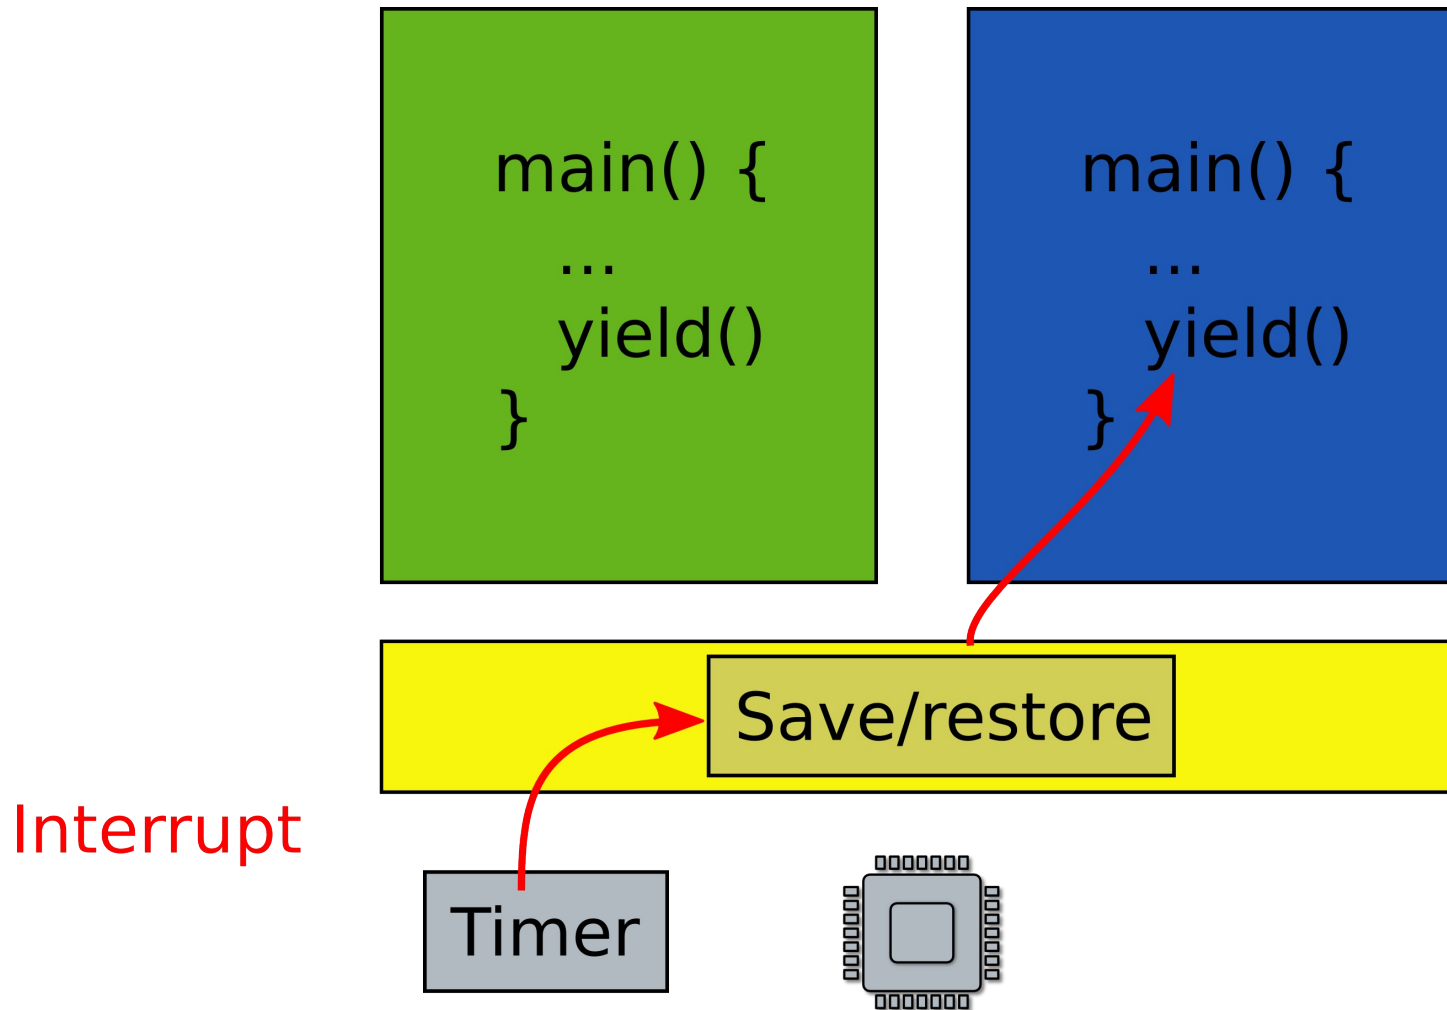


- Linker magic?

Virtual address spaces

- Illusion of a private memory for each application
 - Keep a description of an address space
 - In one of the registers
- All normal program addresses are inside the address space
- OS maintains description of address spaces
 - Switches between them

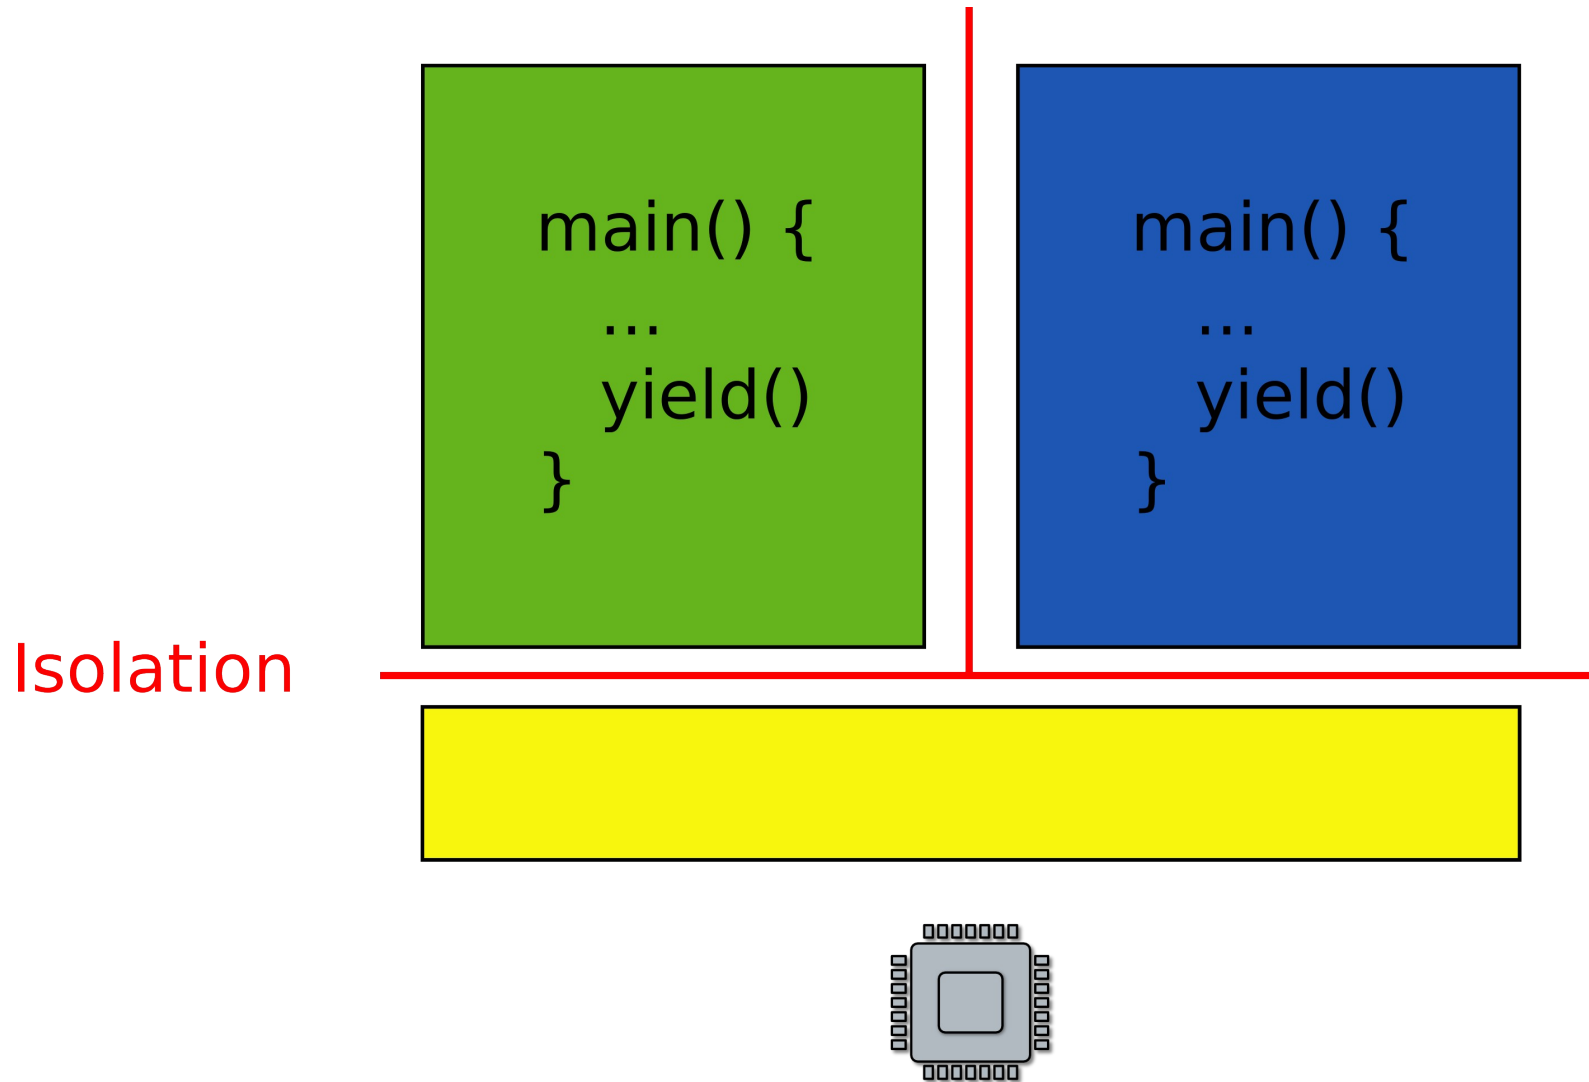
- What if one program fails to release the CPU?
- It will run forever. Need a way to preempt it. How?



Scheduling

- Pick which application to run next
 - And for how long
- Illusion of a private CPU for each task
 - Frequent context switching


- What if one faulty program corrupts the kernel?
- Or other programs?



Isolation

- Today is done with address spaces in hardware
 - Many issues, e.g. shared device drivers, files, etc.
- Can it be done in software?

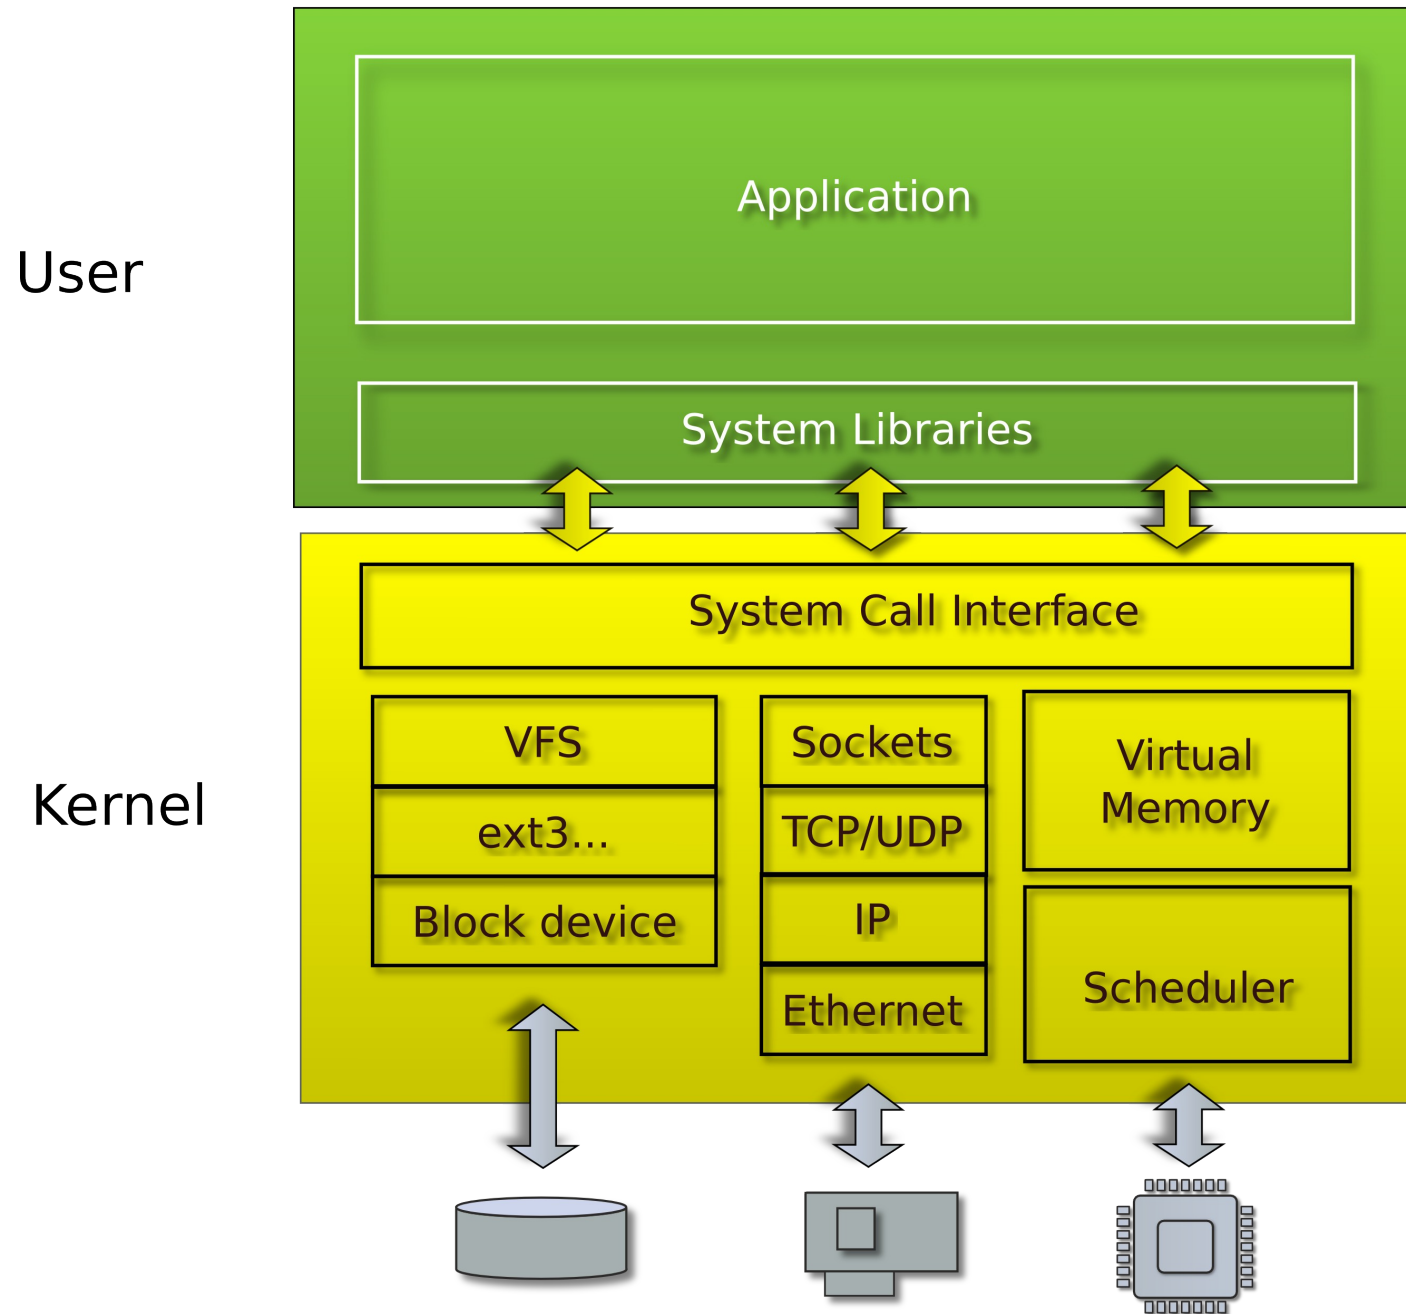
- What about communication?
- Can we invoke a function in a kernel?



- What if you want to save some data?
- Permanent storage
 - E.g., disks
- But disks are just arrays of blocks
 - `wrtie(block_number, block_data)`
- Files
 - High level abstraction for saving data
 - `fd = open("contacts.txt");`
 - `fpritrnf(fd, "Name:%s\n", name);`

- What if you want to send data over the network?
- Network interfaces
 - Send/receive Ethernet packets (Level 2)
 - Two low level
- Sockets
 - High level abstraction for sending data

- Linux/Windows/Mac




Multiple levels of abstraction


- Multiple programs
 - Each has illusion of a private memory and CPU
 - Context switching, scheduling, isolation, communication
- File systems
 - Multiple files, concurrent I/O requests
 - Consistency, caching
- Network protocols
 - Multiple virtual network connections
- Memory management

Virtualization

- Want to run a Windows application on Linux?

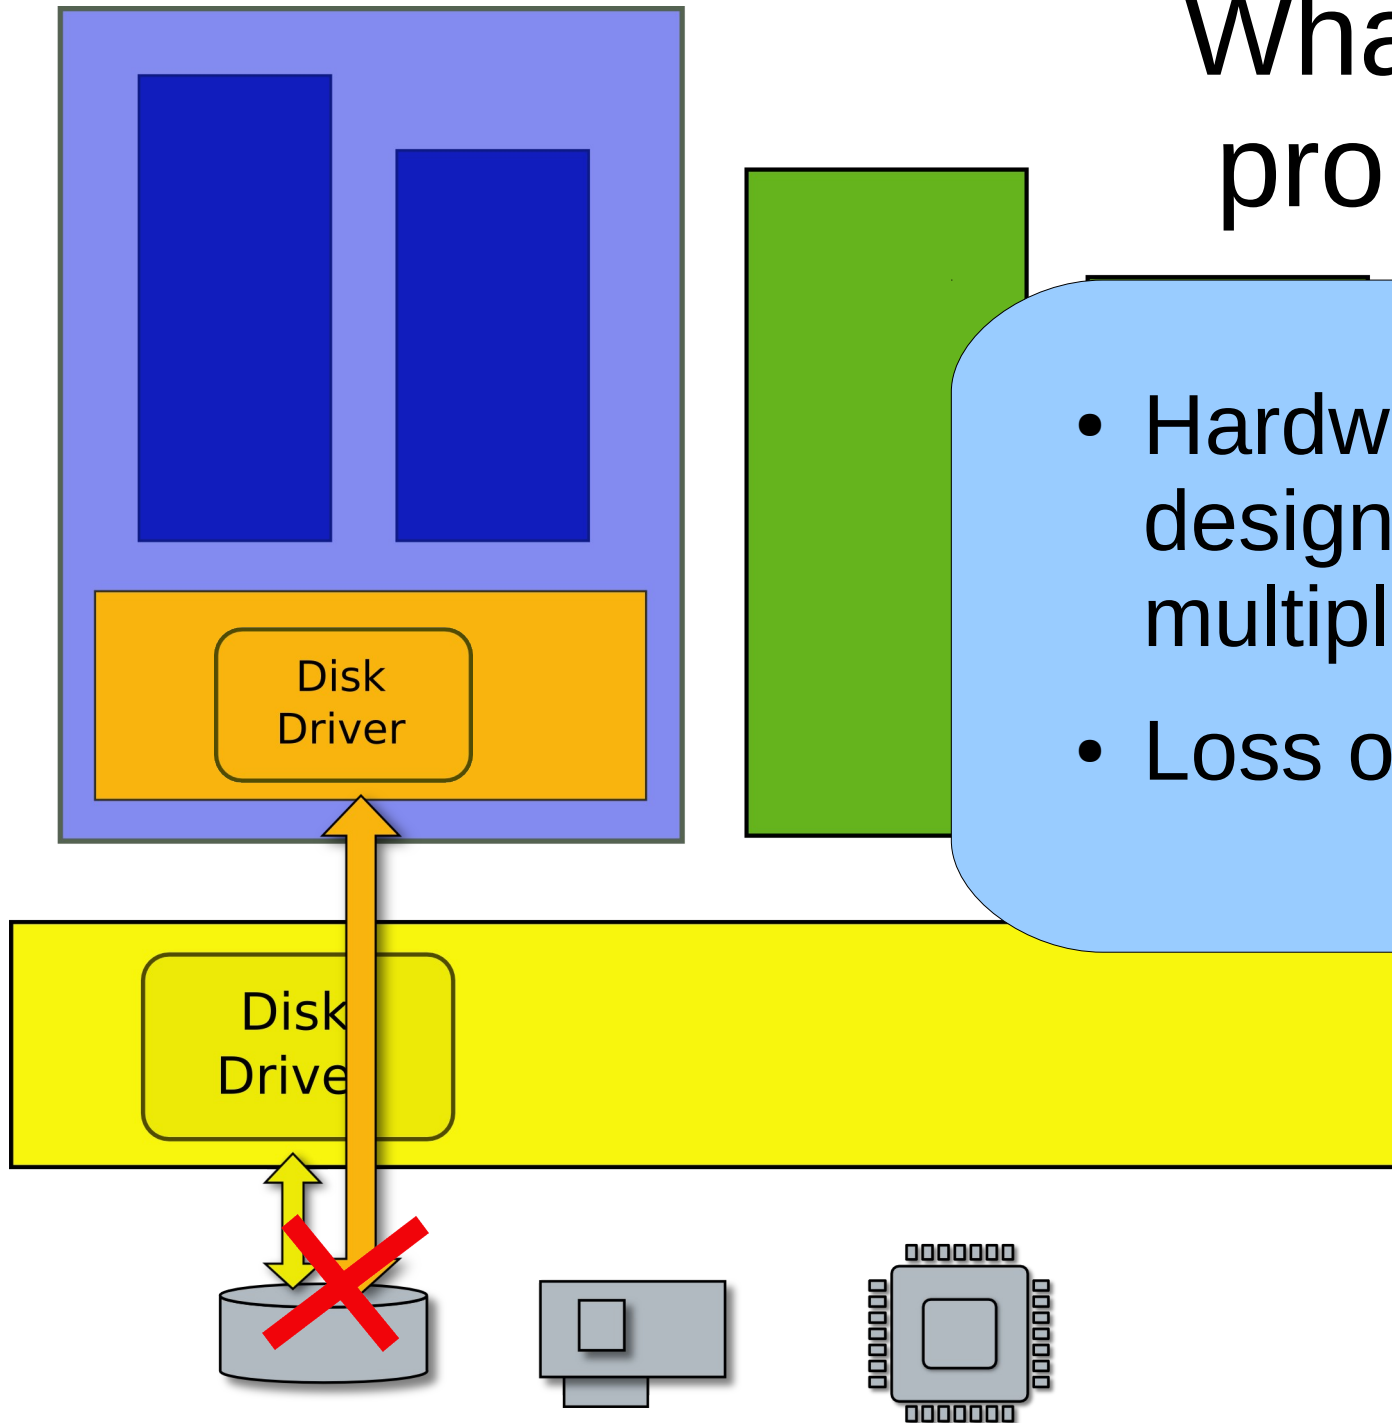


- Want to run a Windows application on Linux?



What is the problem?

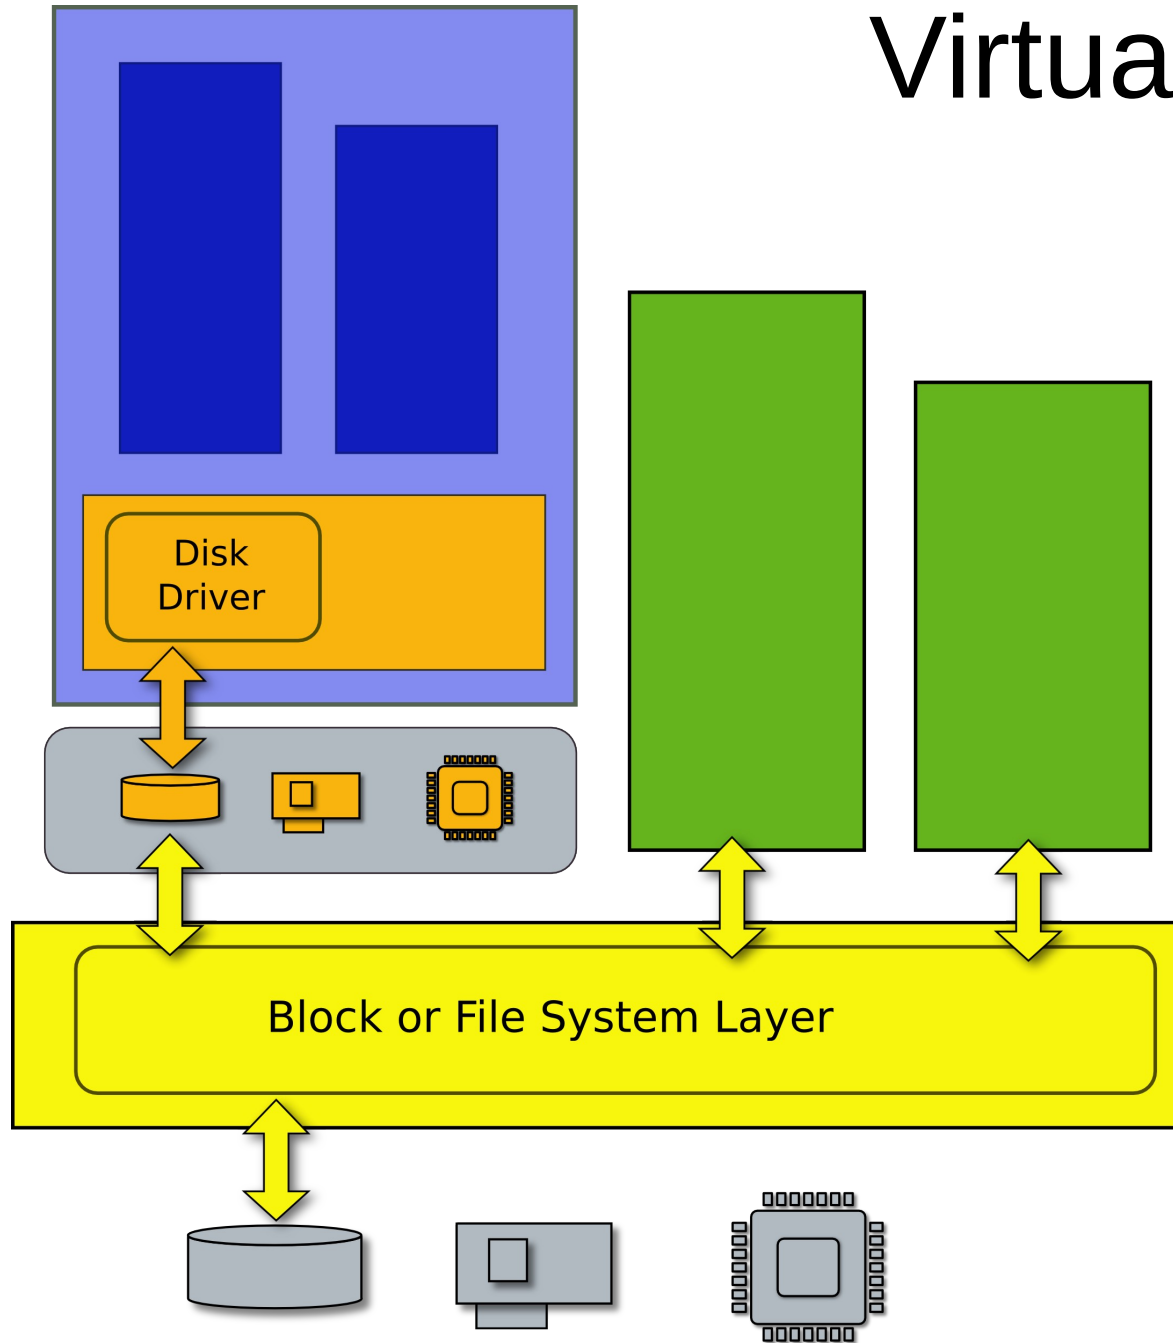
- Hardware is not designed to be multiplexed
- Loss of isolation



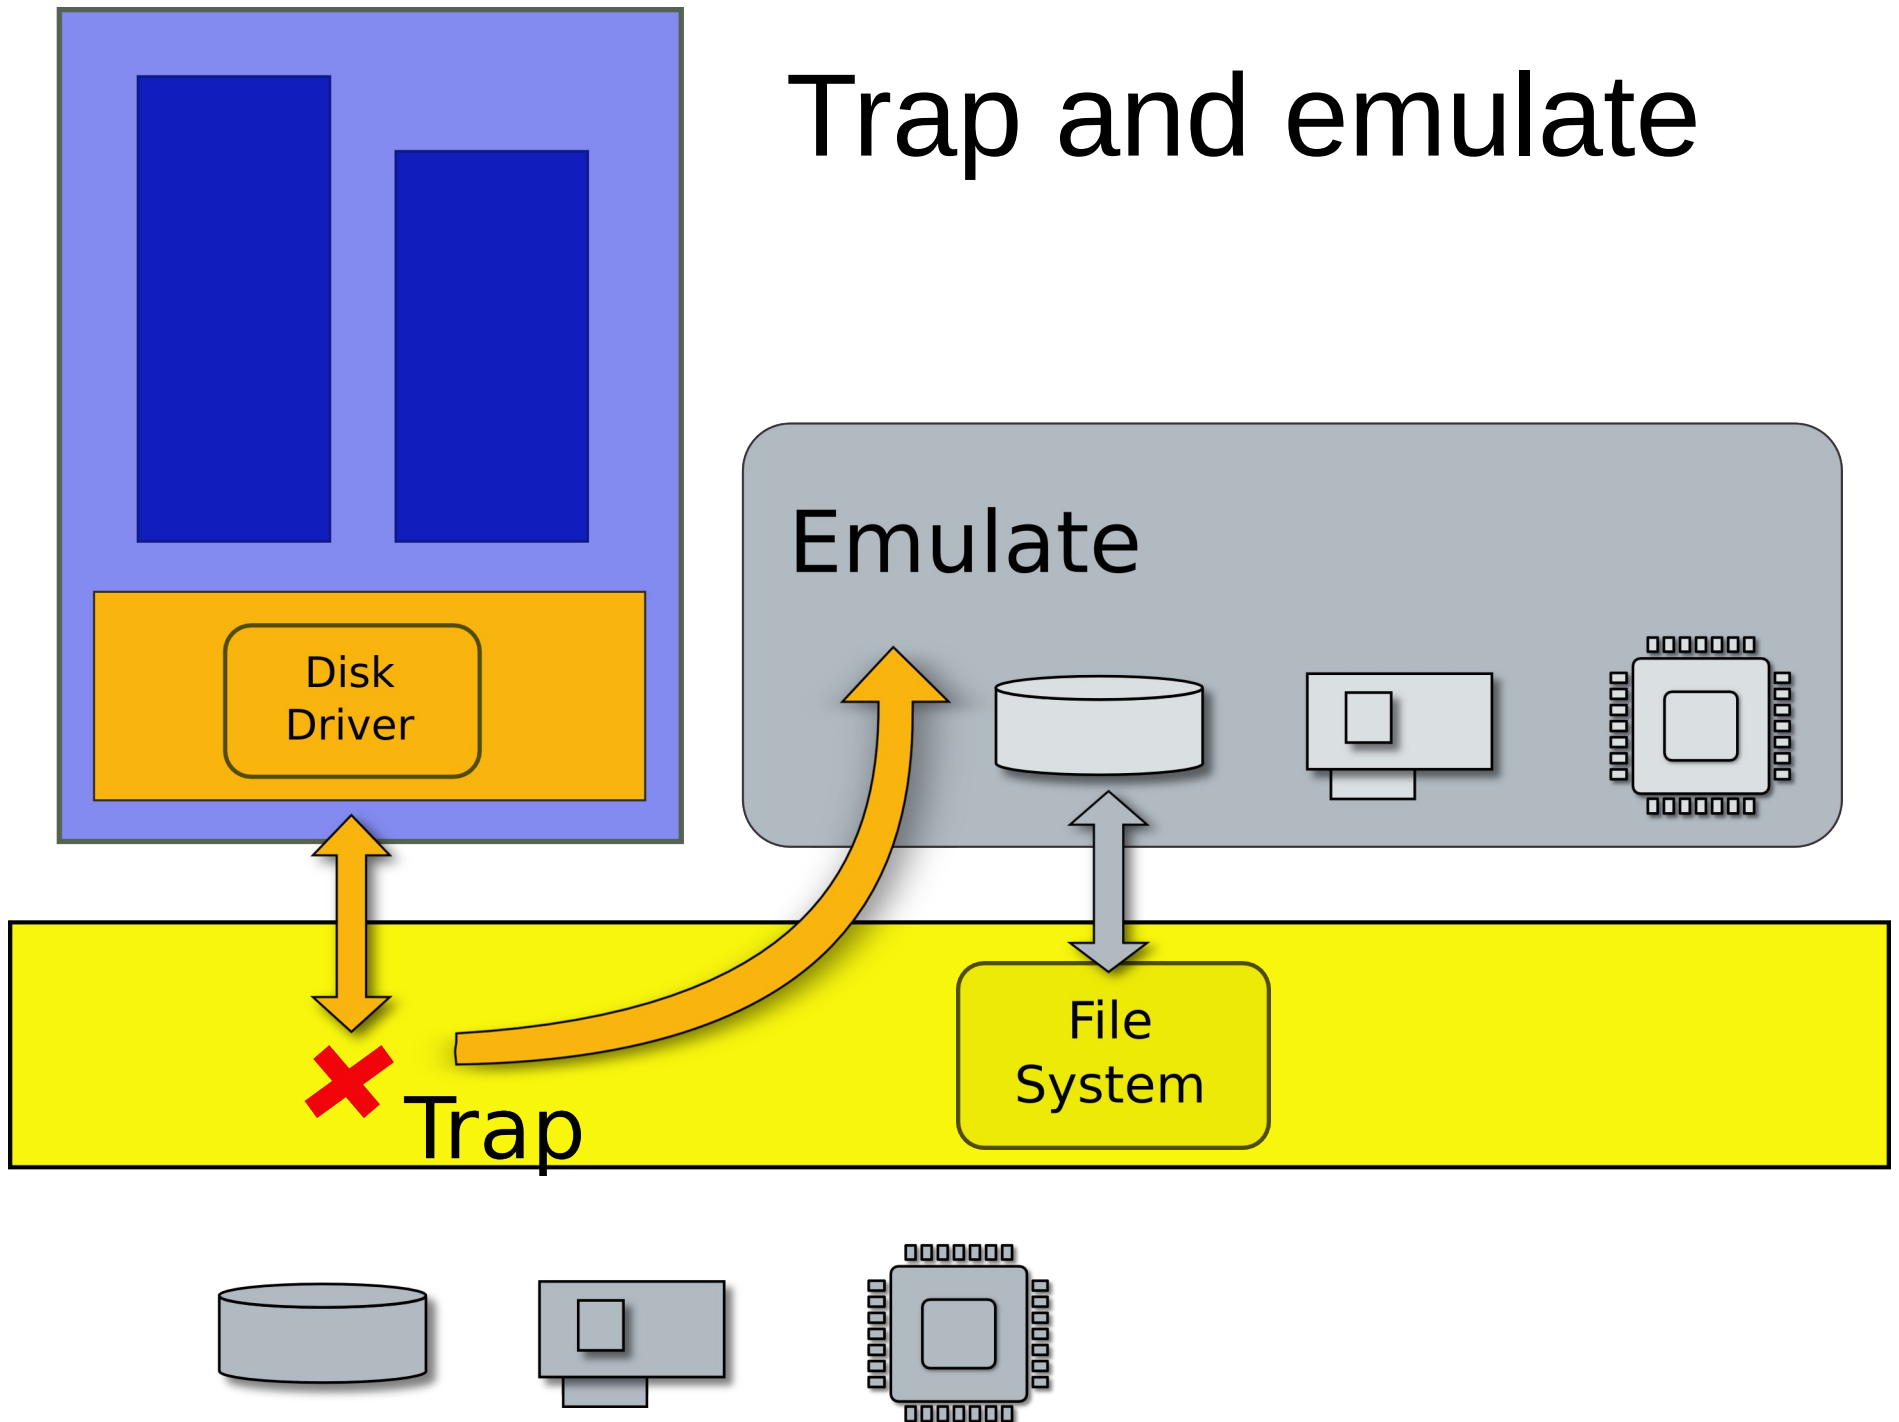
Virtual machine

Efficient duplicate
of a real machine

- Compatibility
- Performance
- Isolation



Trap and emulate



Questions?