

# 143A: Principles of Operating Systems

## Lecture 5: Calling conventions

Anton Burtsev  
January, 2017

# Stack and procedure calls

# Stack

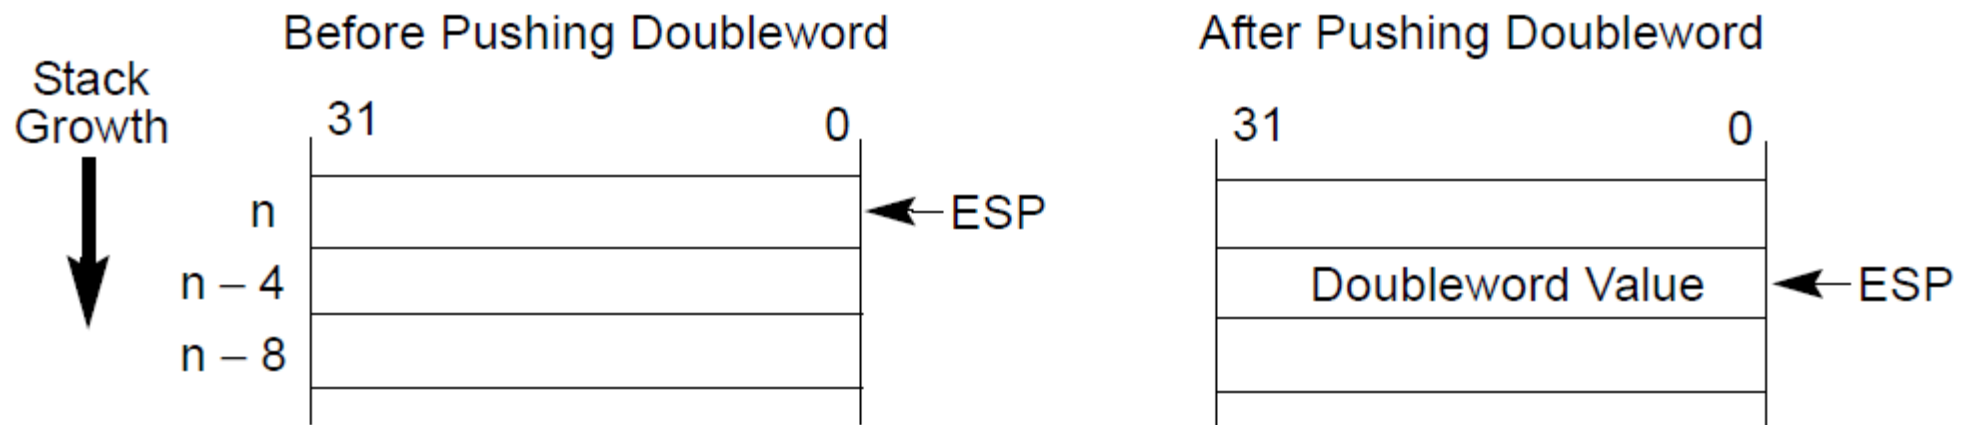
- Main purpose:
  - **Store the return address** for the current procedure
  - Caller pushes return address on the stack
  - Callee pops it and jumps
- Other uses:
  - Local data storage
  - Parameter passing
  - Evaluation stack
    - Register spill

# Manipulating stack

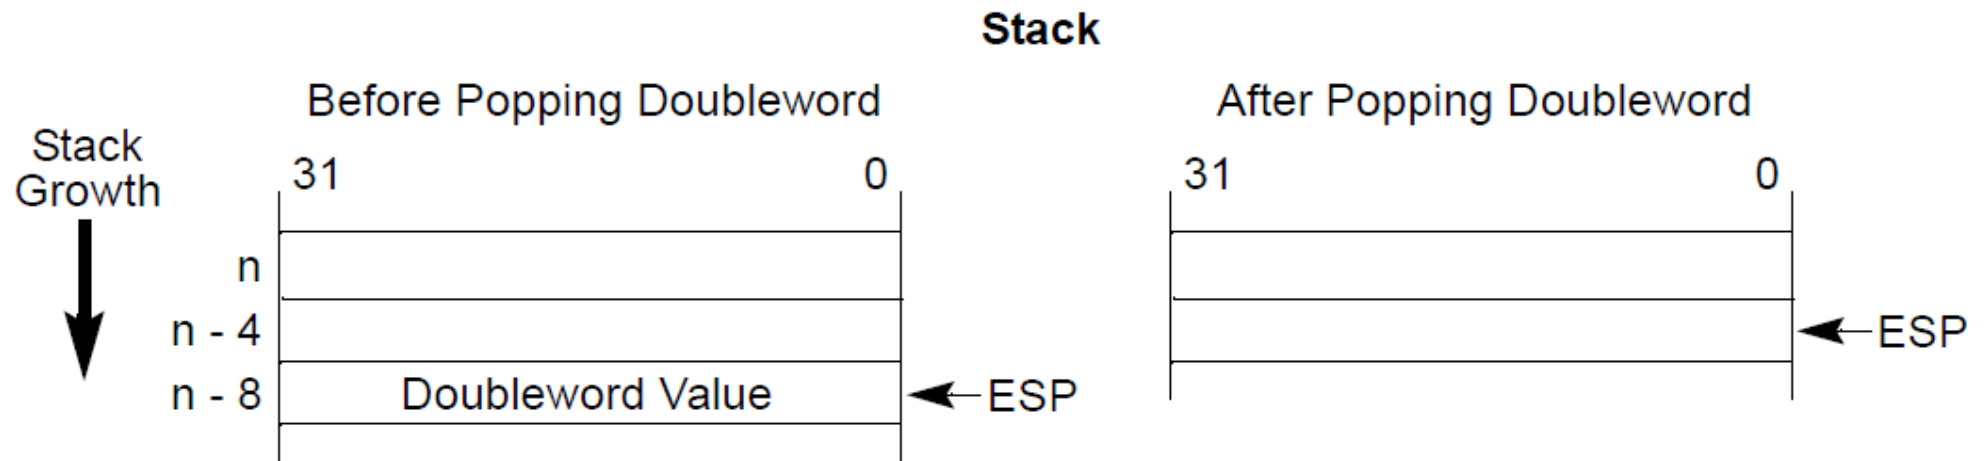
- ESP register
  - Contains the memory address of the topmost element in the stack
- PUSH/POP instructions
  - Insert/remove data on the stack
  - Subtract/add 4 to ESP

# Example: PUSH

## Stack



# Example: POP



# Call/return

- CALL instruction
  - Makes an unconditional jump to a subprogram and pushes the address of the next instruction on the stack

```
push eip + 2 ; save return address
```

```
jmp _my_function
```

- RET instruction
  - Pops off an address and jumps to that address


# Calling conventions

- Goal: reentrant programs
  - How to pass arguments
    - On the stack?
    - In registers?
  - How to return values
    - On the stack?
    - In registers?
- Conventions differ from compiler, optimizations, etc.



# Stack consists of frames

- Each function has a new frame
- Use dedicated register **EBP** (frame pointer)
  - Points to the base of the frame



# Prologue/epilogue

- Each function maintains the frame
  - Uses prologue (blue), and epilogue (yellow)

my\_function:

```
push ebp      ; save original EBP value on stack
mov ebp, esp  ; new EBP = ESP
....          ; function body
pop ebp       ; restore original EBP value
ret
```

# How to allocate local variables

```
void my_function()  
{  
    int a, b, c;  
    ...  
}
```

# Allocating local variables

On the stack!

- Each function has a private instance
- Can call recursively

```
foo(int x) {  
    int a, b, c;  
    a = x + 1;  
    if ( a < 100 )  
        foo(a);  
    return;  
}
```

# Allocating local variables

- Stored right after the saved EBP value in the stack
- Allocated by subtracting the number of bytes required from ESP

`_my_function:`

<code>push ebp</code>	<code>; save original EBP value on stack</code>
<code>mov ebp, esp</code>	<code>; new EBP = ESP</code>
<code>sub esp, LOCAL_BYTES</code>	<code>; = # bytes needed by locals</code>
<code>...</code>	<code>; function body</code>
<code>mov esp, ebp</code>	<code>; deallocate locals</code>
<code>pop ebp</code>	<code>; restore original EBP value</code>
<code>ret</code>	

# Example

```
void my_function() {  
    int a, b, c;  
    ...  
}
```

```
_my_function:  
    push ebp        ; save the value of ebp  
    mov ebp, esp    ; set ebp to be top of the stack (esp)  
    sub esp, 12     ; move esp down to allocate space for the  
                    ; local variables on the stack
```

- With frames local variables can be accessed by dereferencing EBP

```
mov [ebp - 4], 10 ; location of variable a  
mov [ebp - 8], 5  ; location of b  
mov [ebp - 12], 2 ; location of c
```

# How to pass arguments?

- Options
  - Registers
  - On the stack

# How to pass arguments?

- x86 32 bit
  - Pass arguments on the stack
  - Return value is in EAX and EDX
- x86 64 bit – more registers!
  - Pass first 6 arguments in registers
    - RDI, RSI, RDX, RCX, R8, and R9
  - The rest on the stack
  - Return value is in RAX and RDX



# x86\_32: passing arguments on the stack

- Example function

```
void my_function(int x, int y, int z)
{ ... }
```

- Example invocation

```
my_function(2, 5, 10);
```

- Generated code

```
push 10
push 5
push 2
call _my_function
```

# Example stack

:	:	
10	[ebp + 16]	(3rd function argument)
5	[ebp + 12]	(2nd argument)
2	[ebp + 8]	(1st argument)
RA	[ebp + 4]	(return address)
FP	[ebp]	(old ebp value)
	[ebp - 4]	(1st local variable)
:	:	
:	:	
	[ebp - X]	(esp - the current stack pointer)

# Example: callee side code

`_my_function:`

`push ebp`

`mov ebp, esp`

`sub esp, 12 ; allocate local variables`

`; sizeof(a) + sizeof(b) + sizeof(c)`

`; x = [ebp + 8], y = [ebp + 12], z = [ebp + 16]`

`; a=[ebp-4]=[esp+8],`

`; b=[ebp-8]=[esp+4], c=[ebp-12] = [esp]`

`mov esp, ebp ; deallocate local variables`

`pop ebp`

`ret`

```
void my_function(int x, int y, int z)
{
    int a, b, c;
    ...
    return;
}
```

# Example: caller side code

```
int callee(int, int, int);
```

```
int caller(void)
```

```
{
```

```
    int ret;
```

```
    ret = callee(1, 2, 3);
```

```
    ret += 5;
```

```
    return ret;
```

```
}
```

```
caller:
```

```
    ; make new call frame
```

```
    push    ebp
```

```
    mov     ebp, esp
```

```
    ; push call arguments
```

```
    push    3
```

```
    push    2
```

```
    push    1
```

```
    ; call subroutine 'callee'
```

```
    call    callee
```

```
    ; remove arguments from frame
```

```
    add     esp, 12
```

```
    ; use subroutine result
```

```
    add     eax, 5
```

```
    ; restore old call frame
```

```
    pop     ebp
```

```
    ; return
```



```
    ret
```

# Back to stack frames, so why do we need them?

- ... They are not strictly required
- GCC compiler option `-fomit-frame-pointer` can disable them

Don't keep the frame pointer in a register for functions that don't need one. This avoids the instructions to save, set up and restore frame pointers; it also makes an extra register available in many functions. **It also makes debugging impossible on some machines.**

# Stack frames



Initially parameter is

- [ESP + 4]

Later as the function pushes things on the stack it changes, e.g.

- [ESP + 8]

- Debugging becomes hard
  - As ESP changes one has to manually keep track where local variables are relative to ESP (ESP + 4 or +8)
    - Compiler can do this!
    - But it's hard for a human
  - It's hard to unwind the stack in case of a crash
    - To print out a backtrace

# And you only save...

- A couple instructions required to maintain the stack frame
- And 1 register (EBP)
  - x32 has 8 registers (and one is ESP)
    - So taking another one is 12.5% of register space
    - Sometimes its worse it!
  - x64 has 16 registers, so it doesn't really matter
- That said GCC sets `-fomit-frame-pointer` to “on”
  - At -O, -O1, -O2 ...
  - Don't get surprised



# Saving register state across invocations

- Processor doesn't save registers
  - General purpose, segment, flags
- Again, a calling convention is needed
  - Agreement on what gets saved by a callee and caller

# Saving register state across invocations

- Registers EAX, ECX, and EDX are caller-saved
  - The function is free to use them
- ... the rest are callee-saved
  - If the function uses them it has to restore them to the original values

- In general there multiple calling conventions
  - We described **cdecl**
  - **Make sure you know what you're doing**
  - [https://en.wikipedia.org/wiki/X86\\_calling\\_conventions#List\\_of\\_x86\\_calling\\_conventions](https://en.wikipedia.org/wiki/X86_calling_conventions#List_of_x86_calling_conventions)
  - It's easy as long as you know how to read the table

Questions?

# References

- [https://en.wikibooks.org/wiki/X86\\_Disassembly/Functions\\_and\\_Stack\\_Frames](https://en.wikibooks.org/wiki/X86_Disassembly/Functions_and_Stack_Frames)
- [https://en.wikipedia.org/wiki/Calling\\_convention](https://en.wikipedia.org/wiki/Calling_convention)
- [https://en.wikipedia.org/wiki/X86\\_calling\\_conventions](https://en.wikipedia.org/wiki/X86_calling_conventions)
- <http://stackoverflow.com/questions/14666665/trying-to-understand-gcc-option-fomit-frame-pointer>