

# ICS143A: Principles of Operating Systems

Final recap, sample questions

Anton Burtsev  
March, 2017


# What is operating system?


What is a process?


What mechanisms are involved into implementation of a process?

Describe the x86 address translation pipeline (draw figure), explain stages.


# Address translation



























What is the linear address? What address is in the registers, e.g., in %eax?

# Logical and linear addresses

- Segment selector (16 bit) + offset (32 bit)





What segments do the following instructions use? push, jump, mov


What's on the stack? Describe layout of a stack and how it changes during function invocation?

# Example stack

:	:	
10	[ebp + 16]	(3rd function argument)
5	[ebp + 12]	(2nd argument)
2	[ebp + 8]	(1st argument)
RA	[ebp + 4]	(return address)
FP	[ebp]	(old ebp value)
	[ebp - 4]	(1st local variable)
:	:	
:	:	
	[ebp - X]	(esp - the current stack pointer)


Describe the steps and data structures involved into a user to kernel transition (draw diagrams)

# Interrupt path



Which stack is used for execution of an interrupt handler? How does hardware find it?



# Interrupt path



Describe organization of the memory allocator in xv6?



# Physical page allocator



Where did free memory come from?

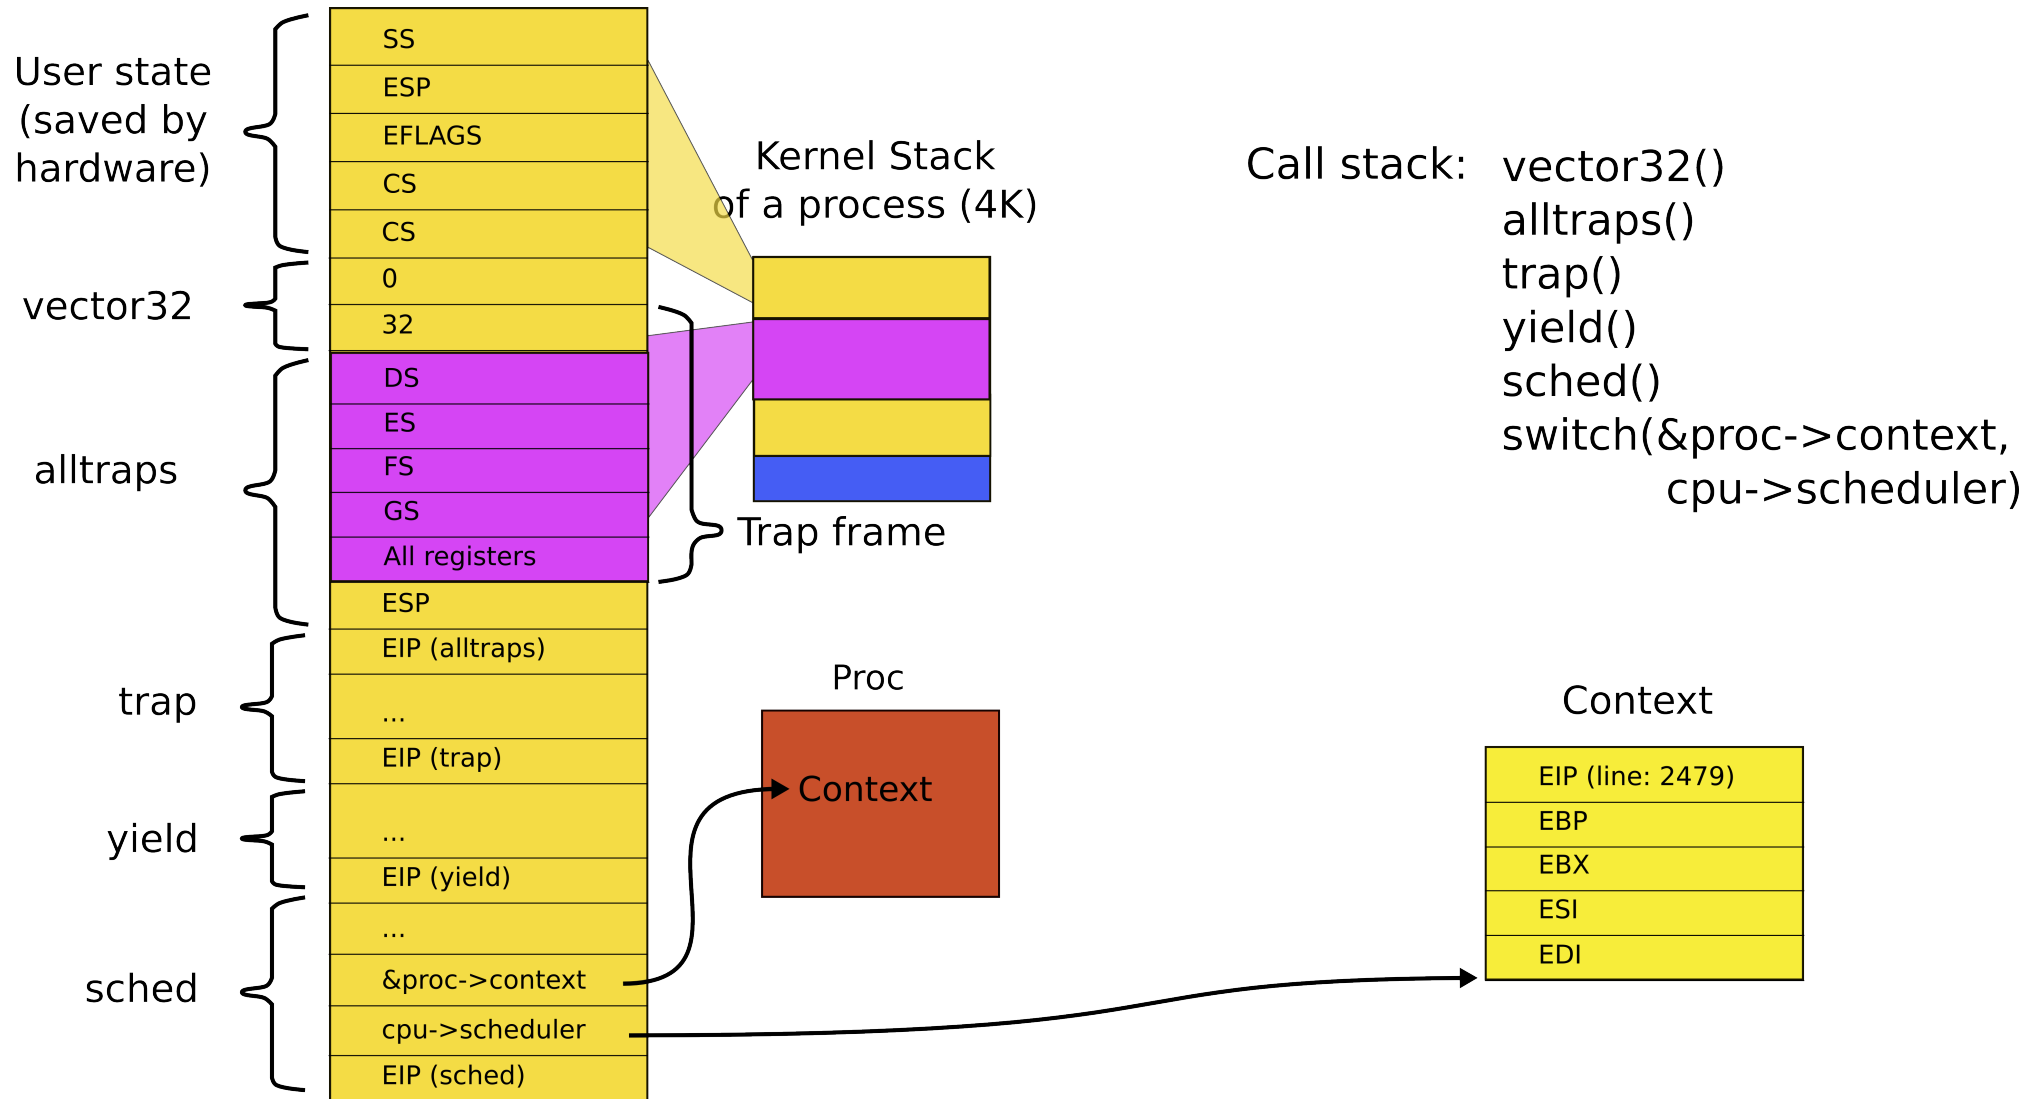
How do we switch between processes?

```
2958 swtch:
2959 movl 4(%esp), %eax
2960 movl 8(%esp), %edx
2961
2962 # Save old callee-save registers
2963 pushl %ebp
2964 pushl %ebx
2965 pushl %esi
2966 pushl %edi
2967
2968 # Switch stacksh
2969 movl %esp, (%eax)
2970 movl %edx, %esp
2971
2972 # Load new callee-save registers
2973 popl %edi
2974 popl %esi
2975 popl %ebx
2976 popl %ebp
2977 ret
```

# swtch()

```
2093 struct context {
2094     uint edi;
2095     uint esi;
2096     uint ebx;
2097     uint ebp;
2098     uint eip;
2099 };
```

# Stack inside switch()



What is the interface between the kernel and user-level processes?


```
3374 void
3375 syscall(void)
3376 {
3377     int num;
3378
3379     num = proc->tf->eax;
3380     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
3381         proc->tf->eax = syscalls[num]();
3382     } else {
3383         cprintf("%d %s: unknown sys call %d\n",
3384             proc->pid, proc->name, num);
3385         proc->tf->eax = -1;
3386     }
3387 }
```

```
3374 void
3375 syscall(void)
3376 {
3377     int num;
3378
3379     num = proc->tf->eax;
3380     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
3381         // proc->tf->eax = syscalls[num]();
3382         proc->tf->esp -= 4;
3383         *(int*)ptoc->tf->esp = syscalls[num]();
3384     } else {
3385         cprintf("%d %s: unknown sys call %d\n",
3386             proc->pid, proc->name, num);
3387         // proc->tf->eax = -1;
3388         proc->tf->esp -= 4;
3389         *(int*)ptoc->tf->esp = -1;
3390     }
3391 }
```




Why do we need locks?

# Request queue (e.g. incoming network packets)




- Linked list, list is pointer to the first element


# CPU 1 and 2 allocate new request



# CPU 1 and 2 update next pointer




# CPU 1 updates head pointer




# CPU2 updates head pointer

list = l



# State after the race



# List implementation with locks

```
9 insert(int data)
10 {
11     struct list *l;
13     l = malloc(sizeof *l);
        acquire(&listlock);
14     l->data = data;
15     l->next = list;
16     list = l;
        release(&listlock);
17 }
```

- Critical section




# Correct implementation

```
1573 void
1574 acquire(struct spinlock *lk)
1575 {
...
1580     // The xchg is atomic.
1581     while(xchg(&lk->locked, 1) != 0)
1582         ;
...
1592 }
```

# Xchg instruction

- Swap a word in memory with a new value
  - Atomic!
  - Return old value


# Deadlocks



# Lock ordering

- Locks need to be acquired in the same order

# Locks and interrupts



# Locks and interrupts


- Never hold a lock with interrupts enabled

# Send/receive queue

```
201 void*
202 send(struct q *q, void *p)
203 {
204     while(q->ptr != 0)
205         ;
206     q->ptr = p;
207     wakeup(q); /*wake recv*/
208 }
```

```
210 void*
211 recv(struct q *q)
212 {
213     void *p;
214
215     while((p = q->ptr) == 0)
216         sleep(q);
217     q->ptr = 0;
218     return p;
219 }
```

# Lost wakeup problem





# The role of file systems

- Sharing
  - Sharing of data across users and applications
- Persistence
  - Data is available after reboot

# Architecture

- On-disk and in-memory data structures represent
  - The tree of named files and directories
  - Record identities of disk blocks which hold data for each file
  - Record which areas of the disk are free

# Crash recovery

- File systems must support crash recovery
  - A power loss may interrupt a sequence of updates
  - Leave file system in inconsistent state
    - E.g. a block both marked free and used

# Multiple users

- Multiple users operate on a file system concurrently
  - File system must maintain invariants

# Speed

- Access to a block device is several orders of magnitude slower
  - Memory: 200 cycles
  - Disk: 20 000 000 cycles
- A file system must maintain a cache of disk blocks in memory

# Block layer

System calls	File descriptors
Pathnames	Recursive lookup
Directories	Directory inodes
Files	Inodes and block allocator
Transactions	Logging
Blocks	Buffer cache

- Read and write data
  - From a block device
  - Into a buffer cache
- Synchronize across multiple readers and writers

# Transactions

- Group multiple writes into an atomic transaction

System calls	File descriptors
Pathnames	Recursive lookup
Directories	Directory inodes
Files	Inodes and block allocator
Transactions	Logging
Blocks	Buffer cache

# Files

- Unnamed files
  - Represented as inodes
  - Sequence of blocks holding file's data

System calls	File descriptors
Pathnames	Recursive lookup
Directories	Directory inodes
Files	Inodes and block allocator
Transactions	Logging
Blocks	Buffer cache



# Directories

System calls	File descriptors
Pathnames	Recursive lookup
Directories	Directory inodes
Files	Inodes and block allocator
Transactions	Logging
Blocks	Buffer cache

- Special kind of inode
  - Sequence of directory entries
  - Each contains name and a pointer to an unnamed inode

# Pathnames

System calls	File descriptors
Pathnames	Recursive lookup
Directories	Directory inodes
Files	Inodes and block allocator
Transactions	Logging
Blocks	Buffer cache


- Hierarchical path names
  - /usr/bin/sh
  - Recursive lookup

# System call

System calls	File descriptors
Pathnames	Recursive lookup
Directories	Directory inodes
Files	Inodes and block allocator
Transactions	Logging
Blocks	Buffer cache


- Abstract UNIX resources as files
  - Files, sockets, devices, pipes, etc.
- Unified programming interface

# File system layout on disk



- Block #0: Boot code
- Block #1: Metadata about the file system
  - Size (number of blocks)
  - Number of data blocks
  - Number of inodes
  - Number of blocks in log

# File system layout on disk



- Block #2 (inode area)
- Bit map area: track which blocks are in use
- Data area: actual file data
- Log area: maintaining consistency in case of a power outage or system crash

```
begin_op();  
...  
bp = bread(...);  
bp->data[...] = ...;  
log_write(bp);  
...  
end_op();
```

# Typical use of transactions

# Strawman scheduler (xv6)

- Organize all processes as a simple list
- In `schedule()`:
  - Pick first one on list to run next
  - Put suspended task at the end of the list
- Problem?

# Xv6 scheduler

```
2458 scheduler(void)
2459 {
2462     for(;;){
2468         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2469             if(p->state != RUNNABLE)
2470                 continue;
2475             proc = p;
2476             switchvm(p);
2477             p->state = RUNNING;
2478             swtch(&cpu->scheduler, proc->context);
2479             switchkvm();
2483             proc = 0;
2484         }
2487     }
2488 }
```



# Strawman scheduler (xv6)

- Organize all processes as a simple list
- In `schedule()`:
  - Pick first one on list to run next
  - Put suspended task at the end of the list
- Problem?

# Strawman scheduler

- Organize all processes as a simple list
- In `schedule()`:
  - Pick first one on list to run next
  - Put suspended task at the end of the list
- Problem?
  - Only allows round-robin scheduling
  - Can't prioritize tasks

# Priority based scheduling

- Higher-priority processes run first
- Processes within the same priority are round-robin

# $O(1)$ scheduler (Linux 2.6 – 2.6.22)


- Priority based scheduling
- Goal: decide who to run next, independent of number of processes in system
  - Still maintain ability to prioritize tasks, handle partially unused quanta, etc

# $O(1)$ data structures


- runqueue: a list of runnable processes
  - Blocked processes are not on any runqueue
  - A runqueue belongs to a specific CPU
  - Each task is on exactly one runqueue
  - Task only scheduled on runqueue's CPU unless migrated
- $2 * 40 * \text{\#CPUs}$  runqueues
  - 40 dynamic priority levels (more later)
  - 2 sets of runqueues – one active and one expired

# $O(1)$ data structures (contd)

Active



Expired




# $O(1)$ intuition


- Take the first task off the lowest-numbered runqueue on active set
  - Confusingly: a lower priority value means higher priority
- When done, put it on appropriate runqueue on expired set
- Once active is completely empty, swap which set of runqueues is active and expired
- Constant time, since fixed number of queues to check; only take first item from non-empty queue

# O(1) example

Active




Expired






# What now?

Active



Flip active and  
expired queues

Expired




# Blocked tasks




- What if a program blocks on I/O, say for the disk?
  - It still has part of its quantum left
  - Not runnable, so don't waste time putting it on the active or expired runqueues
- We need a “wait queue” associated with each blockable event
  - Disk, lock, pipe, network socket, etc.

# Blocking example

Active

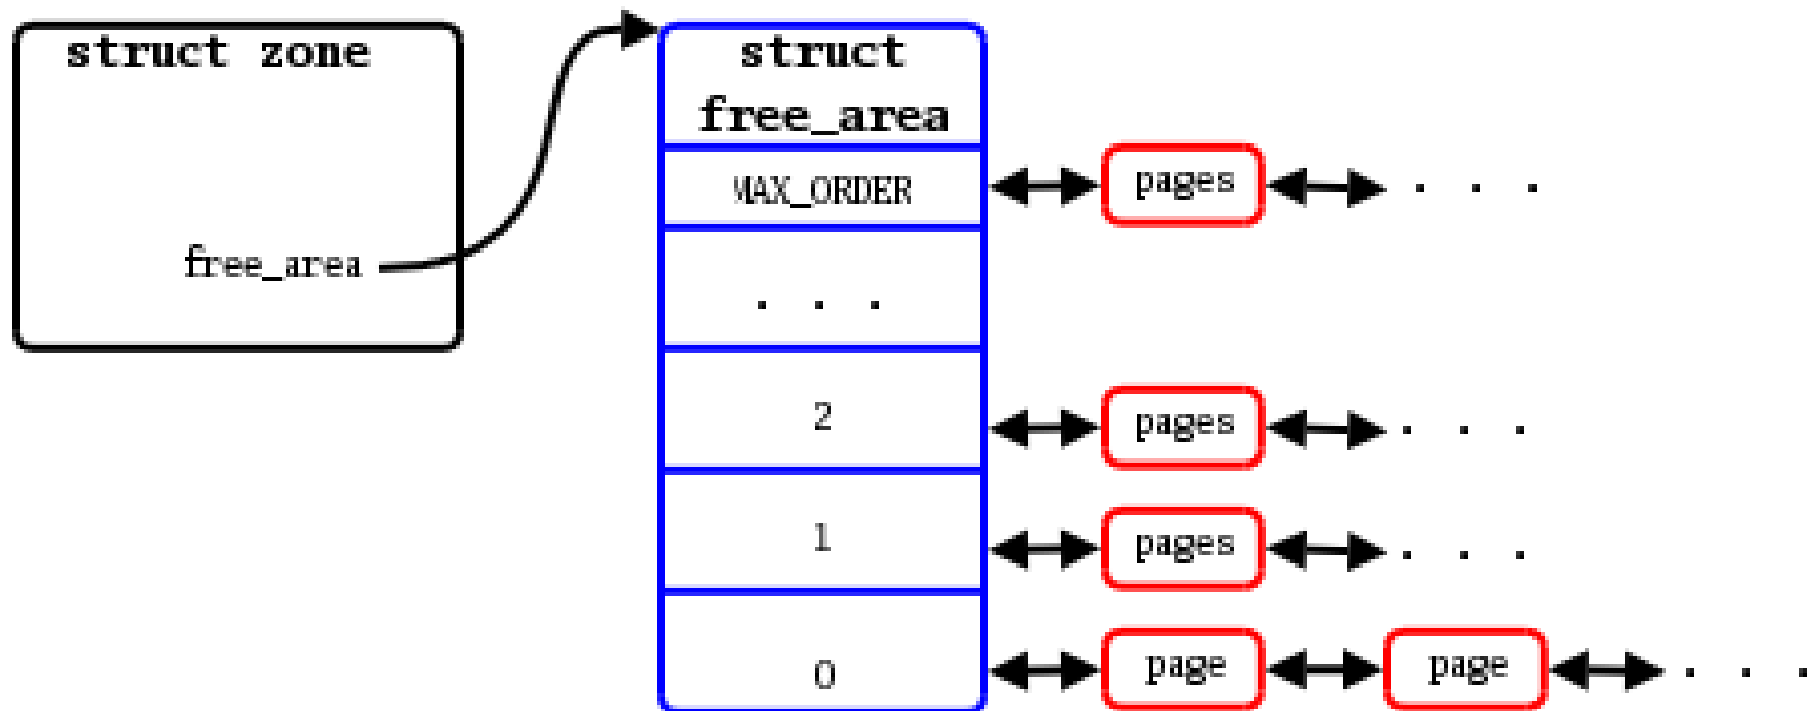


Expired

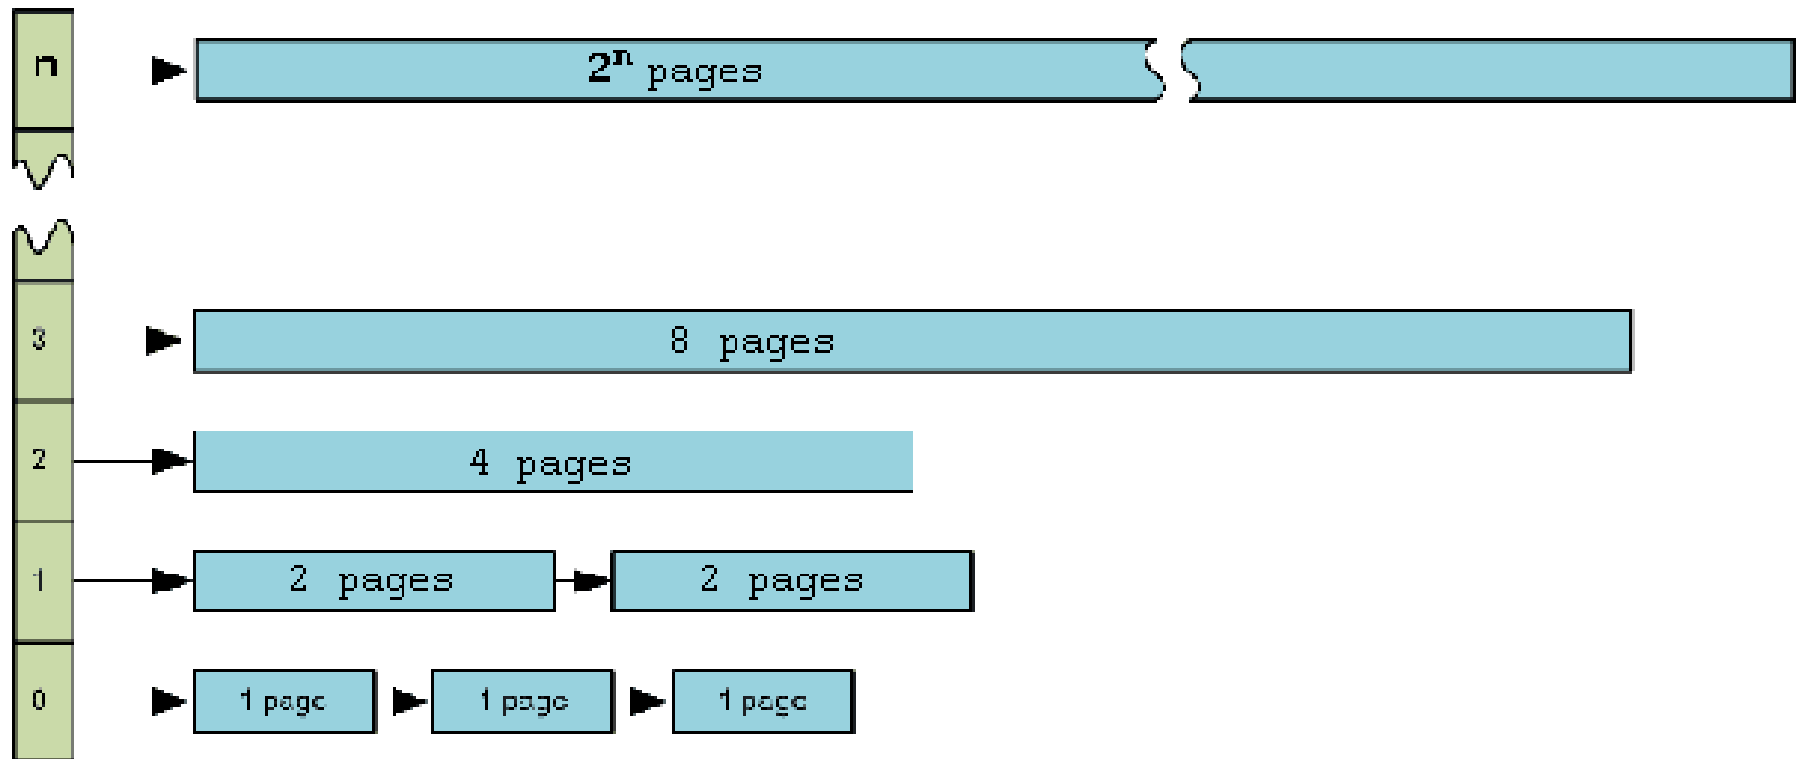


Process goes on  
disk wait queue

# Buddy memory allocator



# Buddy allocator



What's wrong with buddy?

# What's wrong with buddy?


- Buddy allocator is ok for large allocations
  - E.g. 1 page or more
- But what about small allocations?
  - Buddy uses the whole page for a 4 bytes allocation
    - Wasteful
  - Buddy is still slow for short-lived objects

Slab:  
Allocator for object of a fixed size



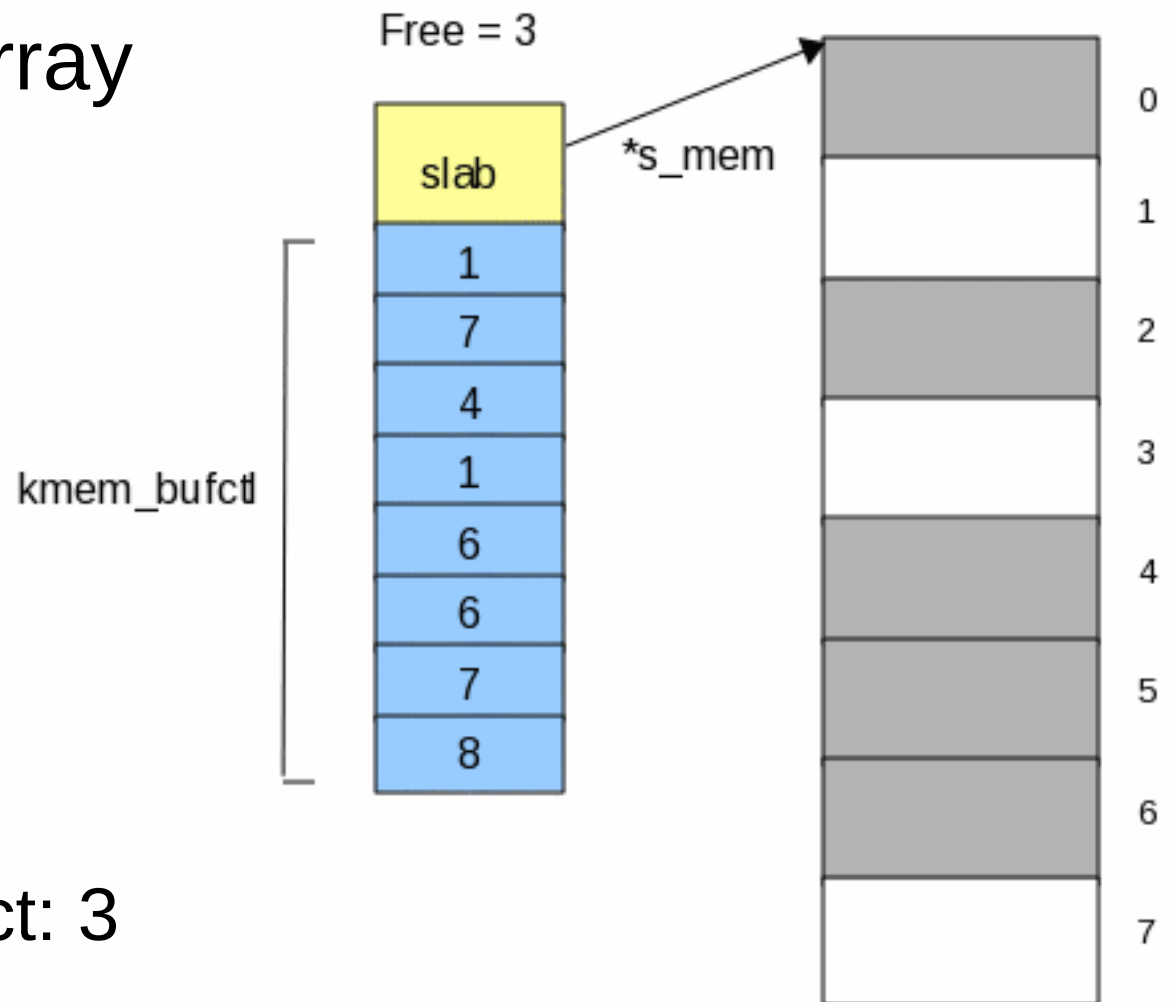
# Slab

- A 2 page slab with 6 objects




# Keeping track of free objects

- kmem\_bufctl array is effectively a linked list



- First free object: 3
- Next free object: 1

# A cache is formed out of slabs




Slab is fine, but what's wrong?

# Slab is fine, but what's wrong?

- We can only allocate objects of one size

# Kmalloc(): variable size objects

- A table of caches
  - Size: 32, 64, 128, etc.



Thank you!