

143A: Principles of Operating Systems

Lecture 10: Address spaces


Anton Burtsev
October, 2017

Recap of the boot sequence

- Setup segments (data and code)
- Switched to protected mode
 - Loaded GDT (segmentation is on)
- Setup stack (to call C functions)
- Loaded kernel from disk
- Setup first page table
 - 2 entries [0 : 4MB] and [2GB : (2GB + 4MB)]
- Setup high-address stack
- Jumped to main()

What's next?

Our goal: 2GB/2GB address space





- Kernel needs normal 2 level page table
 - Right now we have only two entries
 - And current page table is only good for 4MB pages
- But to create page tables we need memory
 - Where can it come from?

Physical page allocator



- Goal:
 - List of free physical pages
 - To allocate page tables, stacks, data structures, etc.
 - Remember current page table is only 1! page
- Where to get memory to keep the list itself?
 - 1 level, only 4MB entries
 - You don't even have space to keep the second level page tables

Physical page allocator





Protected Mode

Physical page allocator





Protected Mode

Physical page allocator



Protected Mode

Physical page allocator



Protected Mode

kalloc() - kernel allocator

```
3087 char*
3088 kalloc(void)
3089 {
3090     struct run *r;
3091     ...
3092     r = kmem.freelist;
3093     if(r)
3094         kmem.freelist = r->next;
3095     ...
3096     return (char*)r;
3097 }
```

```
3065 kfree(char *v)
3066 {
3067     struct run *r;
3068     ...
3077     r = (struct run*)v;
3078     r->next = kmem.freelist;
3079     kmem.freelist = r;
3080     ...
2832 }
```

Kernel needs malloc()

```
1316 int
1317 main(void)
1318 {
1319     kinit1(end, P2V(4*1024*1024)); // phys page allocator
1320     kvmalloc(); // kernel page table
1321     mpinit(); // detect other processors
1322     lapicinit(); // interrupt controller
1323     seginit(); // segment descriptors
1324     cprintf("\ncpu%d: starting xv6\n\n", cpunum());
1325     picinit(); // another interrupt controller
1326     ioapicinit(); // another interrupt controller
1327     consoleinit(); // console hardware
1328     uartinit(); // serial port
...
1340 }
```

```
3030 kinit1(void *vstart, void *vend)
```

```
3031 {
```

```
...
```

```
3034     freerange(vstart, vend);
```

```
3035 }
```

```
3051 freerange(void *vstart, void *vend)
```

```
3052 {
```

```
3053     char *p;
```

```
3054     p = (char*)PGROUNDUP((uint)vstart);
```

```
3055     for(; p + PGSIZE <= (char*)vend; p += PGSIZE)
```

```
3056         kfree(p);
```

```
3057 }
```

Back to kinit1()

Wait! Where do we start?

```
1316 int
1317 main(void)
1318 {
1319     kinit1(end, P2V(4*1024*1024)); // phys page allocator
1320     kvmalloc(); // kernel page table
1321     mpinit(); // detect other processors
```

- What is this **end**?

```
1311 extern char end[];
```

Wait! Where do we start?

```
1316 int
1317 main(void)
1318 {
1319     kinit1(end, P2V(4*1024*1024)); // phys page allocator
1320     kvmalloc(); // kernel page table
1321     mpinit(); // detect other processors
```

- What is this **end**?

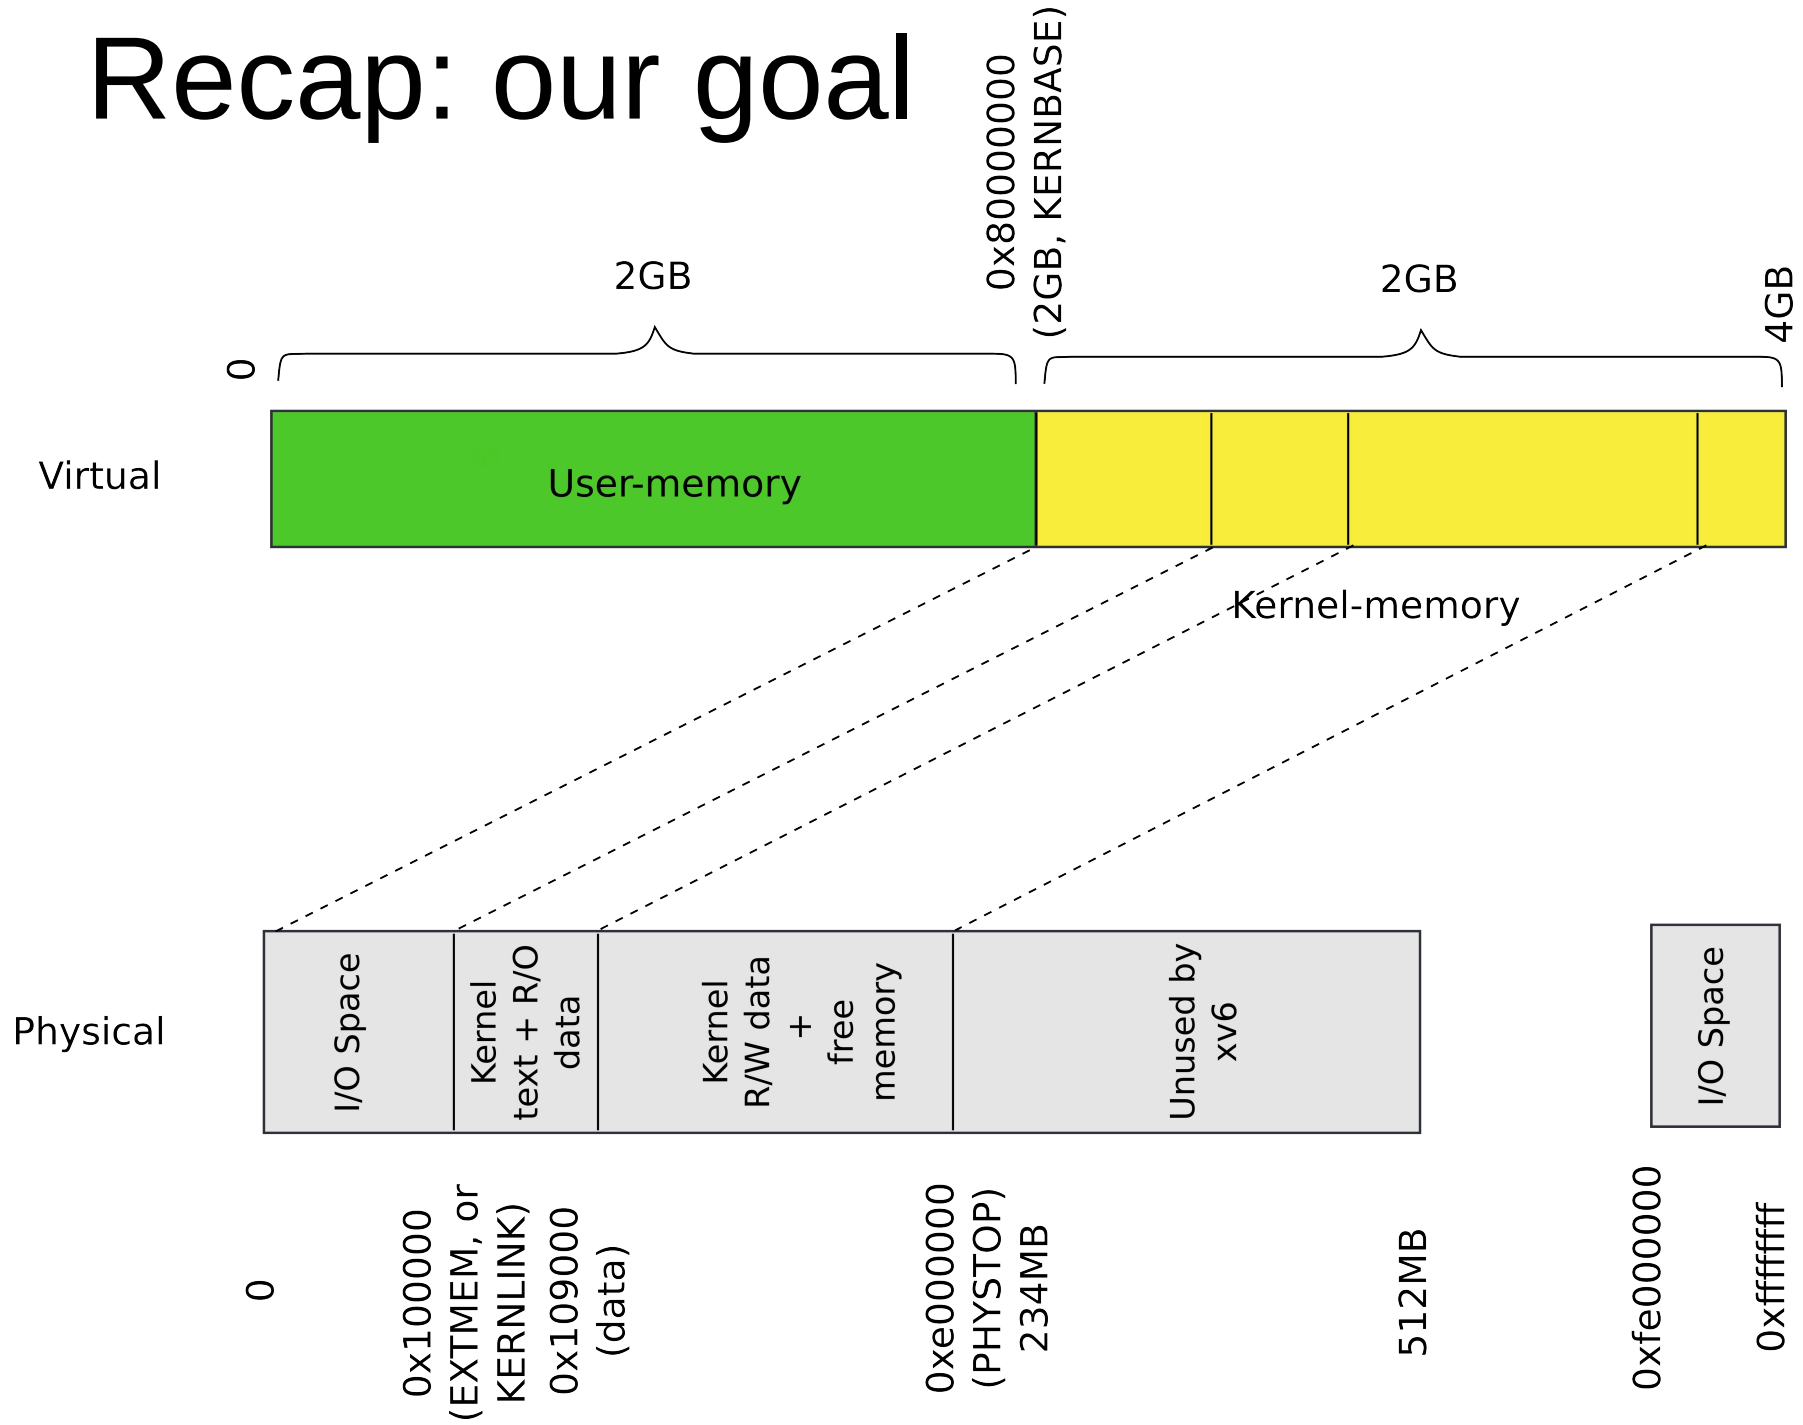
```
1311 extern char end[]; // first address after
                        kernel loaded from ELF file
```


Back to main(): Kernel page table

```
1316 int
1317 main(void)
1318 {
1319     kinit1(end, P2V(4*1024*1024)); // phys page allocator
1320     kvmalloc(); // kernel page table
1321     mpinit(); // detect other processors
1322     lapicinit(); // interrupt controller
1323     seginit(); // segment descriptors
1324     cprintf("\ncpu%d: starting xv6\n\n", cpunum());
1325     picinit(); // another interrupt controller
1326     ioapicinit(); // another interrupt controller
1327     consoleinit(); // console hardware
1328     uartinit(); // serial port
...
1340 }
```

- What do you think has to happen?
 - i.e., how to allocate page tables?

Recap: our goal



Conclusion

- Kernel has a memory allocator
 - It allocates memory in chunks of 4KB
 - Good enough to maintain kernel data structures

Thank you!