

ICS143A: Principles of Operating Systems


Lecture 16: Locking (continued)

Anton Burtsev
February, 2017

Recap: Race conditions

- Disk driver maintains a list of outstanding requests
- Each process can add requests to the list

Request queue (e.g. incoming network packets)



- Linked list, list is pointer to the first element

List implementation with locks

```
9 insert(int data)
10 {
11     struct list *l;
13     l = malloc(sizeof *l);
        acquire(&listlock);
14     l->data = data;
15     l->next = list;
16     list = l;
        release(&listlock);
17 }
```

- Critical section

Xchg instruction

- Swap a word in memory with a new value
 - Atomic!
 - Return old value

Correct implementation

```
1573 void
1574 acquire(struct spinlock *lk)
1575 {
...
1580     // The xchg is atomic.
1581     while(xchg(&lk->locked, 1) != 0)
1582         ;
...
1592 }
```

One last detail...


```
9  insert(int data)
10 {
11     struct list *l;
13     l = malloc(sizeof *l);
        acquire(&listlock);
14     l->data = data;
15     l->next = list;
16     list = l;
        release(&listlock);
17 }
```

Correct implementation

```
1573 void
1574 acquire(struct spinlock *lk)
1575 {
1576     ...
1580     // The xchg is atomic.
1581     while(xchg(&lk->locked, 1) != 0)
1582         ;
1584     // Tell the C compiler and the processor to not move loads or
1585     // stores
1586     // past this point, to ensure that the critical section's memory
1587     // references happen after the lock is acquired.
1588     __sync_synchronize();
1589     ...
1592 }
```


Deadlocks


Deadlocks



Lock ordering

- Locks need to be acquired in the same order

Locks and interrupts



Locks and interrupts

- Never hold a lock with interrupts enabled

```
1573 void
1574 acquire(struct spinlock *lk)
1575 {
1576     pushcli(); // disable interrupts to avoid deadlock.
1577     if(holding(lk))
1578         panic("acquire");
1580     // The xchg is atomic.
1581     while(xchg(&lk->locked, 1) != 0)
1582         ;
1583     ...
1587     __sync_synchronize();
1588     ...
1592 }
```

Disabling interrupts

Simple disable/enable is not enough

- If two locks are acquired
- Interrupts should be re-enabled only after the second lock is released
- Pushcli() uses a counter

```
1655 pushcli(void)
1656 {
1657     int eflags;
1658
1659     eflags = readeflags();
1660     cli();
1661     if(cpu->ncli == 0)
1662         cpu->intena = eflags & FL_IF;
1663     cpu->ncli += 1;
1664 }
...
1667 popcli(void)
1668 {
1669     if(readeflags() & FL_IF)
1670         panic("popcli - interruptible");
1671     if(--cpu->ncli < 0)
1672         panic("popcli");
1673     if(cpu->ncli == 0 && cpu->intena)
1674         sti();
1675 }
```

Pushcli()/popcli()

Locks and interprocess communication

Send/receive queue

```
100 struct q {  
101     void *ptr;  
102 };  
103  
104 void*  
105 send(struct q *q, void *p)  
106 {  
107     while(q->ptr != 0)  
108         ;  
109     q->ptr = p;  
110 }
```

```
112 void*  
113 recv(struct q *q)  
114 {  
115     void *p;  
116  
117     while((p = q->ptr) == 0)  
118         ;  
119     q->ptr = 0;  
120     return p;  
121 }
```

Send/receive queue

```
100 struct q {
101     void *ptr;
102 };
103
104 void*
105 send(struct q *q, void *p)
106 {
107     while(q->ptr != 0)
108         ;
109     q->ptr = p;
110 }
112 void*
113 recv(struct q *q)
114 {
115     void *p;
116
117     while((p = q->ptr) == 0)
118         ;
119     q->ptr = 0;
120     return p;
121 }
```

- Expensive if communication is rare
 - Receiver wastes CPU cycles

Sleep and wakeup syscalls

- `sleep(channel)`
 - Put calling process to sleep
 - Release CPU for other work
- `wakeup(channel)`
 - Wakes all processes sleeping on a channel
 - If any
 - i.e., causes `sleep()` calls to return

Send/receive queue

```
201 void*
202 send(struct q *q, void *p)
203 {
204     while(q->ptr != 0)
205         ;
206     q->ptr = p;
207     wakeup(q); /*wake recv*/
208 }
```

```
210 void*
211 recv(struct q *q)
212 {
213     void *p;
214
215     while((p = q->ptr) == 0)
216         sleep(q);
217     q->ptr = 0;
218     return p;
219 }
```


Send/receive queue

```
201 void*
202 send(struct q *q, void *p)
203 {
204     while(q->ptr != 0)
205         ;
206     q->ptr = p;
207     wakeup(q); /*wake recv*/
208 }
```

```
210 void*
211 recv(struct q *q)
212 {
213     void *p;
214
215     while((p = q->ptr) == 0)
216         sleep(q);
217     q->ptr = 0;
218     return p;
219 }
```

- `recv()` gives up the CPU to other processes
 - But there is a problem...

Lost wakeup problem



Thank you!